

# Open-Source Project

Deployment Guide

## HCLS AI Factory

Deployment and Configuration Guide  
for NVIDIA DGX Spark

*Open-Source Precision Medicine Platform  
on NVIDIA DGX Spark*

02/2026 | Version 1.0 | Apache 2.0 License

Author: Adam Jones

# Table of Contents

- [\*\*1. Introduction\*\*](#)
- [\*\*2. Architecture Overview\*\*](#)
- [\*\*3. Prerequisites\*\*](#)
- [\*\*4. Environment Preparation\*\*](#)
- [\*\*5. Repository Setup\*\*](#)
- [\*\*6. Reference Data Preparation\*\*](#)
- [\*\*7. Docker Compose Configuration\*\*](#)
- [\*\*8. Deploy Genomics Pipeline \(Stage 1\)\*\*](#)
- [\*\*9. Deploy RAG Chat Pipeline \(Stage 2\)\*\*](#)
- [\*\*10. Deploy Drug Discovery Pipeline \(Stage 3\)\*\*](#)
- [\*\*11. Nextflow Orchestration\*\*](#)
- [\*\*12. Service Startup and Health\*\*](#)
- [\*\*13. Monitoring and Observability\*\*](#)
- [\*\*14. Security Configuration\*\*](#)
- [\*\*15. Data Management\*\*](#)
- [\*\*16. Performance Tuning\*\*](#)
- [\*\*17. Troubleshooting Guide\*\*](#)
- [\*\*18. VCP/FTD Demo Walkthrough\*\*](#)
- [\*\*19. Scaling Beyond DGX Spark\*\*](#)
- [\*\*20. Appendix A: Complete Configuration Reference\*\*](#)
- [\*\*21. Appendix B: API Reference\*\*](#)
- [\*\*22. Appendix C: Schema Definitions\*\*](#)
- [\*\*23. Appendix D: Docker Image Reference\*\*](#)
- [\*\*24. Appendix E: Validation Checklists\*\*](#)
- [\*\*25. Appendix F: Glossary\*\*](#)

# 1. Introduction

## 1.1 Purpose

This document provides step-by-step instructions for deploying the HCLS AI Factory on an NVIDIA DGX Spark workstation. It covers all three pipeline stages — genomics, RAG-powered variant intelligence, and AI-driven drug discovery — using exclusively open-source and publicly available components.

## 1.2 Scope

The guide addresses hardware validation, software installation, container deployment, data preparation, pipeline execution, monitoring, security, and troubleshooting. It targets the open-source fork of the HCLS AI Factory that runs entirely on Docker Compose without requiring VAST Data, Kubernetes, or multi-node infrastructure.

## 1.3 Audience

- **Bioinformatics Engineers** deploying genomics pipelines on DGX Spark
- **ML/AI Engineers** integrating RAG and BioNeMo NIM microservices
- **DevOps Engineers** managing containerized service stacks
- **Researchers** forking the project for their own precision medicine workflows

## 1.4 Document Conventions

Convention	Meaning
<code>monospace</code>	Commands, file paths, code
<b>Bold</b>	UI elements, key terms
<i>Italic</i>	Variable values to be replaced
<code>\$VARIABLE</code>	Environment variable
<code>&lt;placeholder&gt;</code>	User-supplied value

## 1.5 Genomics and Drug Discovery Primer

This section provides essential background for engineers who may not have a biology or chemistry background.

### 1.5.1 DNA Sequencing

DNA sequencing reads the order of nucleotide bases (A, T, C, G) in an organism's genome. Modern short-read sequencers (e.g., Illumina) produce paired-end reads — two sequences from opposite ends of a DNA fragment. The standard demo sample HG002 is a 30x whole-genome sequencing (WGS) dataset with 2x250 bp paired-end reads, producing approximately 200 GB of FASTQ data.

### 1.5.2 Genomics Pipeline Stages

Stage	Input	Tool	Output	Description
Quality Control	FASTQ	FastQC	QC Report	Assess read quality and adapter contamination
Alignment	FASTQ + Reference	BWA-MEM2 (fq2bam)	BAM	Map reads to GRCh38 reference genome
Variant Calling	BAM	DeepVariant	VCF	Identify SNPs and indels vs. reference
Annotation	VCF	VEP + ClinVar + AlphaMissense	Annotated VCF	Add functional, clinical, and pathogenicity data
Embedding	Annotated VCF	BGE-small-en-v1.5	Vectors (384-dim)	Convert variant evidence to dense embeddings

### 1.5.3 Variant Annotation

Variants are annotated from multiple sources:

- **VEP (Variant Effect Predictor):** Assigns functional consequences and impact levels — HIGH, MODERATE, LOW, or MODIFIER.
- **ClinVar:** NCBI database of 4.1 million clinical variant interpretations (Pathogenic, Likely Pathogenic, Benign, etc.).
- **AlphaMissense:** DeepMind model with 71,697,560 missense variant pathogenicity predictions. Thresholds: pathogenic (>0.564), ambiguous (0.34-0.564), benign (<0.34).

### 1.5.4 Vector Embeddings and RAG

Annotated variants are converted to 384-dimensional dense vectors using the BGE-small-en-v1.5 embedding model and stored in Milvus. Retrieval-Augmented Generation (RAG) queries Milvus for relevant genomic evidence, then passes the results as context to Anthropic Claude for natural-language clinical interpretation.

### 1.5.5 Drug Discovery Pipeline

The 10-stage drug discovery pipeline transforms a genomic target into ranked drug candidates:

Stage	Name	Description
1	Initialize	Load configuration, validate target gene and variant
2	Normalize Target	Map gene symbol to UniProt ID and canonical name
3	Structure Discovery	Query RCSB PDB for 3D protein structures, score by resolution and method
4	Structure Preparation	Download PDB files, extract binding site coordinates
5	Molecule Generation	Generate SMILES candidates via MolMIM NIM (Port 8001) using seed molecule
6	Chemistry QC	Filter by Lipinski Rule of Five (MW<=500,

		LogP<=5, HBD<=5, HBA<=10)
7	Conformer Generation	Generate 3D conformers with RDKit for docking input
8	Molecular Docking	Score binding affinity via DiffDock NIM (Port 8002)
9	Composite Ranking	Rank candidates: 30% generation + 40% docking + 30% QED
10	Reporting	Generate PDF report with structures, scores, and recommendations

### 1.5.6 End-to-End Data Flow Summary

```

FASTQ (200 GB) → Parabricks fq2bam → BAM (100 GB) → DeepVariant → VCF (11.7M variants)
→ Annotation (ClinVar + AlphaMissense + VEP) → Milvus (384-dim vectors)
→ Claude RAG (variant interpretation) → Target Hypothesis
→ PDB Structure Retrieval → MolMIM (molecule generation)
→ DiffDock (molecular docking) → Composite Ranking → PDF Report

```

## 2. Architecture Overview

### 2.1 System Components

The HCLS AI Factory comprises three application pipeline stages running on a single DGX Spark:

Stage	Name	Function
Stage 1	Genomics Pipeline	FASTQ alignment and variant calling with GPU-accelerated Parabricks
Stage 2	RAG Chat Pipeline	Variant annotation, vector embedding, and Claude-powered conversational AI
Stage 3	Drug Discovery Pipeline	Structure-aware molecule generation, docking, and composite ranking

### 2.2 Technology Stack

Layer	Technology	Version / Details
Hardware	NVIDIA DGX Spark	GB10 GPU, 128 GB unified LPDDR5x, 144 ARM64 cores
OS	DGX OS	Ubuntu-based, ARM64 (aarch64)
Container Runtime	Docker + NVIDIA Container Toolkit	nvidia-docker runtime
Orchestration	Docker Compose	Multi-service deployment
Pipeline Orchestration	Nextflow	DSL2, multiple profiles
GPU Genomics	NVIDIA Parabricks	4.6.0-1

Vector Database	Milvus	2.4 (with etcd + MinIO)
Embedding Model	BGE-small-en-v1.5	384 dimensions
LLM	Anthropic Claude	claude-sonnet-4-20250514
Molecule Generation	BioNeMo MolMIM NIM	1.0
Molecular Docking	BioNeMo DiffDock NIM	1.0
Cheminformatics	RDKit	Python library
Monitoring	Grafana + Prometheus	10.2.2 / v2.48.0
GPU Monitoring	DCGM Exporter	Port 9400
Language	Python	3.10+

## 2.3 Service Architecture

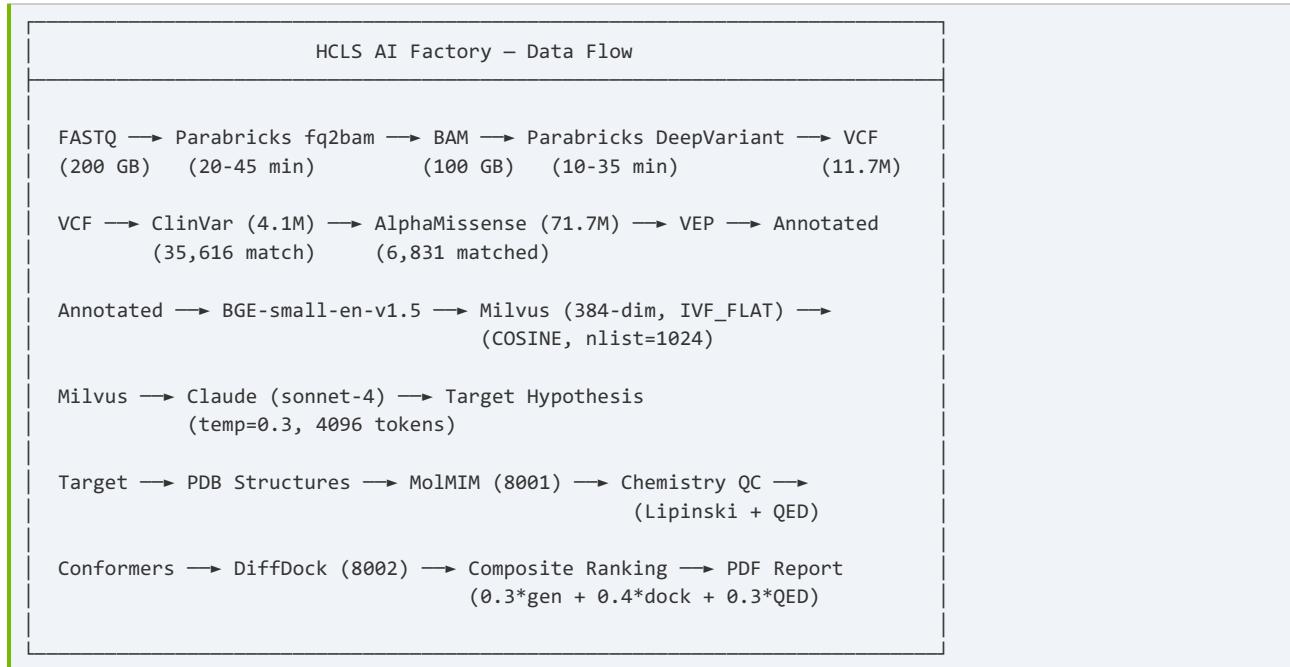
The platform deploys 14 services across 14 ports:

#	Service	Port	Protocol	Description
1	Landing Page	8080	HTTP	Platform entry point and service directory
2	Genomics Portal	5000	HTTP	Genomics pipeline UI and results viewer
3	RAG API	5001	HTTP	REST API for variant queries and RAG
4	Milvus	19530	gRPC	Vector database for genomic evidence
5	Attu	8000	HTTP	Milvus administration UI
6	Streamlit Chat	8501	HTTP	Conversational AI interface for variant analysis
7	MolMIM NIM	8001	HTTP	BioNeMo molecule generation microservice
8	DiffDock NIM	8002	HTTP	BioNeMo molecular docking microservice
9	Discovery UI	8505	HTTP	Drug discovery pipeline interface
10	Discovery Portal	8510	HTTP	Drug discovery results and reporting portal
11	Grafana	3000	HTTP	Monitoring dashboards
12	Prometheus	9099	HTTP	Metrics collection and storage
13	Node Exporter	9100	HTTP	Host system metrics
14	DCGM Exporter	9400	HTTP	NVIDIA GPU metrics

Infrastructure services (not externally exposed):

Service	Port	Purpose
etcd	2379	Milvus metadata store
MinIO	9000	Milvus object storage

## 2.4 Data Flow



## 3. Prerequisites

### 3.1 Hardware Requirements

Component	Specification
System	NVIDIA DGX Spark
GPU	GB10 Grace Blackwell Superchip
Memory	128 GB unified LPDDR5x
CPU	144 ARM64 cores
Architecture	aarch64 (ARM64)
Price	\$3,999

Storage requirements:

Dataset / Component	Size
GRCh38 Reference Genome	3.1 GB

FASTQ Input (HG002 30x WGS)	~200 GB
BAM Output (intermediate)	~100 GB
ClinVar Database	~1.2 GB
AlphaMissense Predictions	~4 GB
Milvus Index Data	~2 GB
BioNeMo Model Cache	~10 GB
<b>Total Minimum</b>	<b>~320 GB</b>
<b>Recommended</b>	<b>1 TB NVMe</b>

## 3.2 Software Requirements

Software	Minimum Version	Notes
DGX OS	Latest	Ubuntu-based ARM64
Docker Engine	24.0+	With Compose V2
NVIDIA Container Toolkit	Latest	nvidia-docker runtime
CUDA Toolkit	12.x	Included with DGX OS
Python	3.10+	For pipeline scripts
Nextflow	23.04+	DSL2 support required
Git	2.30+	For repository clone
NGC CLI	Latest	For BioNeMo container pulls

## 3.3 Network Requirements

- Internet access for initial setup (container pulls, data downloads)
- Outbound HTTPS to [api.anthropic.com](https://api.anthropic.com) for Claude API calls
- Outbound HTTPS to [nvcr.io](https://nvcr.io) for NGC container registry
- Outbound HTTPS to NCBI, RCSB PDB for reference data downloads
- All service ports (listed in Section 2.3) accessible on localhost

## 3.4 Access Credentials

Credential	Purpose	How to Obtain
<a href="#">ANTHROPIC_API_KEY</a>	Claude API access	<a href="https://console.anthropic.com">https://console.anthropic.com</a>
<a href="#">NGC_API_KEY</a>	NVIDIA NGC container registry	<a href="https://ngc.nvidia.com">https://ngc.nvidia.com</a>

# 4. Environment Preparation

## 4.1 DGX Spark Initial Setup

Verify the system is a DGX Spark with the expected hardware:

[BASH](#)

```
# Verify ARM64 architecture
uname -m
# Expected: aarch64

# Verify CPU cores
nproc
# Expected: 144

# Verify total memory (128 GB)
free -h | grep Mem
# Expected: ~128 GB total

# Verify GPU is detected
nvidia-smi
# Expected: GB10 GPU listed with driver version
```

## 4.2 NVIDIA Driver and CUDA Verification

### BASH

```
# Check NVIDIA driver version
nvidia-smi --query-gpu=driver_version --format=csv,noheader
# Expected: 550.x or later

# Check CUDA version
nvcc --version
# Expected: CUDA 12.x

# Verify GPU compute capability
nvidia-smi --query-gpu=compute_cap --format=csv,noheader

# Run a quick GPU test
nvidia-smi -q | head -30
```

## 4.3 Docker Installation and Configuration

### BASH

```
# Verify Docker is installed
docker --version
# Expected: Docker version 24.0+

# Verify Docker Compose V2
docker compose version
# Expected: Docker Compose version v2.x

# Verify NVIDIA runtime is available
docker info | grep -i runtime
# Expected: nvidia runtime listed

# Test GPU access from a container
docker run --rm --gpus all nvidia/cuda:12.4.0-base-ubuntu22.04 nvidia-smi
```

Configure Docker daemon for NVIDIA runtime as default:

**BASH**

```
sudo tee /etc/docker/daemon.json <<'EOF'
{
  "default-runtime": "nvidia",
  "runtimes": {
    "nvidia": {
      "path": "nvidia-container-runtime",
      "runtimeArgs": []
    }
  },
  "default-address-pools": [
    {"base": "172.20.0.0/16", "size": 24}
  ],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "50m",
    "max-file": "3"
  }
}
EOF

sudo systemctl restart docker
```

## 4.4 Python Environment Setup

**BASH**

```
# Verify Python version
python3 --version
# Expected: Python 3.10+

# Create virtual environment
python3 -m venv ~/hcls-env
source ~/hcls-env/bin/activate

# Install core dependencies
pip install --upgrade pip
pip install \
  anthropic \
  pymilvus \
  sentence-transformers \
  rdkit-pypi \
  pydantic \
  streamlit \
  fastapi \
  uvicorn \
  requests \
  pandas \
  numpy \
  reportlab \
  biopython \
  nextflow
```

## 4.5 NGC CLI Installation

**BASH**

```
# Download NGC CLI for ARM64
```

```
wget -O ngc-cli.zip https://api.ngc.nvidia.com/v2/resources/nvidia/ngc-apps/ngc_cli/versions/latest/files/ngccli_arm64.zip

# Extract and install
unzip ngc-cli.zip -d ~/ngc-cli
chmod +x ~/ngc-cli/ngc-cli/ngc
export PATH=$PATH:~/ngc-cli/ngc-cli

# Configure NGC CLI
ngc config set
# Enter your NGC API key when prompted

# Verify authentication
ngc registry image list --format_type csv | head -5
```

## 5. Repository Setup

## 5.1 Fork and Clone

## BASH

```
# Fork the repository on GitHub, then clone your fork
git clone https://github.com/<your-username>/hcls-ai-factory.git
cd hcls-ai-factory

# Verify repository structure
ls -la
```

## 5.2 Repository Layout

```
hcls-ai-factory/
├── docker-compose.yml
├── .env.example
├── nextflow.config
├── main.nf
├── start-services.sh
└── requirements.txt

└── genomics/
    ├── parabricks/
    │   ├── fq2bam.sh
    │   └── deepvariant.sh
    ├── portal/
    │   └── app.py
    └── data/
        ├── reference/
        ├── fastq/
        ├── bam/
        └── vcf/

└── rag/
    ├── api/
    │   └── app.py

# All 14 services + infrastructure
# Template environment configuration
# Nextflow pipeline configuration
# Nextflow DSL2 pipeline definition
# Service startup script
# Python dependencies

# Stage 1: Genomics Pipeline
# Parabricks configs and scripts
# BWA-MEM2 alignment wrapper
# DeepVariant variant calling wrapper
# Genomics Portal (Port 5000)

# Input/output data directory
# GRCh38 reference genome
# Input FASTQ files
# Alignment output
# Variant call output

# Stage 2: RAG Chat Pipeline
# RAG API (Port 5001)
```

```

chat/                               # Streamlit Chat (Port 8501)
|   └── app.py
embeddings/                         # BGE embedding pipeline
|   └── embed_variants.py
annotation/                          # Variant annotation pipeline
|   ├── clinvar.py
|   ├── alphamissense.py
|   └── vep.py
knowledge/                           # Gene knowledge base
|   └── genes.json                  # 201 genes, 13 therapeutic areas
data/
|   ├── clinvar/                   # ClinVar database
|   └── alphamissense/             # AlphaMissense predictions

discovery/                           # Stage 3: Drug Discovery Pipeline
|   ├── pipeline/                  # 10-stage discovery pipeline
|   |   ├── __init__.py
|   |   ├── initialize.py          # Stage 1: Initialize
|   |   ├── normalize.py           # Stage 2: Normalize Target
|   |   ├── structure_discovery.py # Stage 3: Structure Discovery
|   |   ├── structure_prep.py      # Stage 4: Structure Preparation
|   |   ├── molecule_gen.py        # Stage 5: Molecule Generation
|   |   ├── chemistry_qc.py        # Stage 6: Chemistry QC
|   |   ├── conformer_gen.py       # Stage 7: Conformer Generation
|   |   ├── docking.py             # Stage 8: Molecular Docking
|   |   ├── ranking.py              # Stage 9: Composite Ranking
|   |   └── reporting.py           # Stage 10: Reporting
|   ├── ui/                         # Discovery UI (Port 8505)
|   |   └── app.py
|   ├── portal/                     # Discovery Portal (Port 8510)
|   |   └── app.py
|   └── models/                     # Pydantic data models
|       └── schemas.py

monitoring/                          # Monitoring stack
|   ├── grafana/
|   |   ├── provisioning/
|   |   |   └── dashboards/
|   |   └── prometheus/
|   |       └── prometheus.yml
|   └── exporters/

landing/                             # Landing Page (Port 8080)
|   └── index.html

scripts/                             # Utility scripts
|   ├── run_pipeline.py             # Pipeline launcher
|   ├── download_references.sh     # Reference data downloader
|   └── validate_deployment.sh     # Deployment validator

docs/                                # Documentation
|   └── ...

```

## 5.3 Environment Configuration

### BASH

```
# Copy the example environment file
cp .env.example .env
```

```
# Edit with your credentials and paths
nano .env
```

The `.env` file should contain:

**BASH**

```
# === API Keys ===
ANTHROPIC_API_KEY=sk-ant-api03-XXXXXXXXXXXX
NGC_API_KEY=XXXXXXXXXXXX

# === Model Configuration ===
CLAUDE_MODEL=claude-sonnet-4-20250514
CLAUDE_TEMPERATURE=0.3

# === Reference Data ===
REFERENCE_GENOME=/data/reference/GRCh38.fa

# === Milvus Configuration ===
MILVUS_HOST=localhost
MILVUS_PORT=19530

# === BioNeMo NIM URLs ===
MOLMIM_URL=http://localhost:8001
DIFFDOCK_URL=http://localhost:8002

# === Pipeline Configuration ===
PIPELINE_MODE=full
NUM_CANDIDATES=100
MIN_QED=0.67
MIN_DOCK_SCORE=-6.0

# === Monitoring ===
GRAFANA_USER=admin
GRAFANA_PASSWORD=changeme
```

## 5.4 Directory Structure for Data

**BASH**

```
# Create data directories
mkdir -p genomics/data/{reference,fastq,bam,vcf}
mkdir -p rag/data/{clinvar,alphamissense}
mkdir -p discovery/data/{structures,molecules,reports}
mkdir -p monitoring/data/{grafana,prometheus}
```

## 6. Reference Data Preparation

### 6.1 GRCh38 Reference Genome

#### BASH

```
# Download GRCh38 reference genome (~3.1 GB)
cd genomics/data/reference

wget
https://ftp.ncbi.nlm.nih.gov/genomes/all/GCA/000/001/405/GCA_000001405.15_GRCh38/seqs_for_alignment_pipelines
.ucsc_ids/GCA_000001405.15_GRCh38_no_alt_analysis_set.fna.gz

# Decompress
gunzip GCA_000001405.15_GRCh38_no_alt_analysis_set.fna.gz
mv GCA_000001405.15_GRCh38_no_alt_analysis_set.fna GRCh38.fa

# Index the reference (required by Parabricks)
# Note: Parabricks fq2bam can build its own index, but pre-building saves time
samtools faidx GRCh38.fa

# Verify
ls -lh GRCh38.fa*
# Expected: GRCh38.fa (~3.1 GB), GRCh38.fa.fai
```

### 6.2 ClinVar Database

#### BASH

```
# Download ClinVar VCF (~1.2 GB)
cd rag/data/clinvar

wget https://ftp.ncbi.nlm.nih.gov/pub/clinvar/vcf_GRCh38/clinvar.vcf.gz
wget https://ftp.ncbi.nlm.nih.gov/pub/clinvar/vcf_GRCh38/clinvar.vcf.gz.tbi

# Verify record count (~4.1M clinical variants)
zcat clinvar.vcf.gz | grep -v '^#' | wc -l
# Expected: ~4,100,000

echo "ClinVar download complete"
```

### 6.3 AlphaMissense Database

#### BASH

```
# Download AlphaMissense predictions (~4 GB)
cd rag/data/alphamissense

wget https://storage.googleapis.com/dm_alphaMissense/AlphaMissense_hg38.tsv.gz

# Verify record count (~71.7M predictions)
zcat AlphaMissense_hg38.tsv.gz | tail -n +5 | wc -l
# Expected: ~71,697,560
```

```
echo "AlphaMissense download complete"
```

#### AlphaMissense pathogenicity thresholds:

Classification	Score Range	Description
Pathogenic	> 0.564	Likely damaging to protein function
Ambiguous	0.34 - 0.564	Uncertain significance
Benign	< 0.34	Likely tolerated

## 6.4 HG002 Sample Data

### BASH

```
# Download HG002 FASTQ files for demo/testing (~200 GB)
cd genomics/data/fastq

# GIAB HG002 30x WGS, 2x250 bp paired-end
# Note: These are large files – ensure ~200 GB free space
wget ftp://ftp-
trace.ncbi.nlm.nih.gov/ReferenceSamples/giab/data/AshkenazimTrio/HG002_NA24385_son/NIST_HiSeq_HG002_Homogeneity-10953946/NHGRI_Illumina300X_AJtrio_novoalign_bams/HG002.GRCh38.2x250.fastq.gz

# For a smaller test subset, use a downsampled version if available
echo "HG002 download complete – verify file sizes match expected ~200 GB"
ls -lh *.fastq.gz
```

## 7. Docker Compose Configuration

### 7.1 Service Definition Overview

The `docker-compose.yml` defines all 14 application services plus 2 infrastructure services (etcd, MinIO) for Milvus. Services are organized into three groups matching the pipeline stages, plus monitoring.

### 7.2 docker-compose.yml Structure

#### YAML

```
version: '3.8'

services:
  # — Infrastructure —
  etcd:
    image: quay.io/coreos/etcd:v3.5.5
    environment:
      - ETCD_AUTO_COMPACTION_MODE=revision
      - ETCD_AUTO_COMPACTION_RETENTION=1000
    ports:
      - "2379:2379"
    volumes:
```

```

        - etcd_data:/etcd
    restart: unless-stopped

minio:
  image: minio/minio:latest
  environment:
    MINIO_ACCESS_KEY: minioadmin
    MINIO_SECRET_KEY: minioadmin
  ports:
    - "9000:9000"
  volumes:
    - minio_data:/data
  command: server /data
  restart: unless-stopped

# — Milvus Vector Database —
milvus:
  image: milvusdb/milvus:v2.4-latest
  ports:
    - "19530:19530"
  environment:
    ETCD_ENDPOINTS: etcd:2379
    MINIO_ADDRESS: minio:9000
  depends_on:
    - etcd
    - minio
  volumes:
    - milvus_data:/var/lib/milvus
  restart: unless-stopped

attu:
  image: zilliz/attu:latest
  ports:
    - "8000:3000"
  environment:
    MILVUS_URL: milvus:19530
  depends_on:
    - milvus
  restart: unless-stopped

# — Stage 1: Genomics —
genomics-portal:
  build: ./genomics/portal
  ports:
    - "5000:5000"
  volumes:
    - ./genomics/data:/data
  environment:
    - REFERENCE_GENOME=/data/reference/GRCh38.fa
  restart: unless-stopped

# — Stage 2: RAG Chat —
rag-api:
  build: ./rag/api
  ports:
    - "5001:5001"
  environment:
    - ANTHROPIC_API_KEY=${ANTHROPIC_API_KEY}
    - CLAUDE_MODEL=${CLAUDE_MODEL}
    - CLAUDE_TEMPERATURE=${CLAUDE_TEMPERATURE}

```

```

- MILVUS_HOST=milvus
- MILVUS_PORT=19530
depends_on:
... (121 more lines)

```

## 7.3 Infrastructure Services

Milvus 2.4 requires two backend services:

Service	Image	Port	Purpose
etcd	quay.io/coreos/etcd:v3.5.5	2379	Metadata storage for Milvus
MinIO	minio/minio:latest	9000	Object storage for Milvus segments
Milvus	milvusdb/milvus:v2.4-latest	19530	Vector database

## 7.4 Volume Mounts and Data Paths

Volume	Container Path	Host Purpose
./genomics/data	/data	Reference genome, FASTQ, BAM, VCF
./rag/data	/data	ClinVar, AlphaMissense databases
etcd_data	/etcd	Milvus metadata persistence
minio_data	/data	Milvus segment persistence
milvus_data	/var/lib/milvus	Milvus index persistence
prometheus_data	/prometheus	Prometheus TSDB
grafana_data	/var/lib/grafana	Grafana state and dashboards

## 7.5 GPU Resource Allocation

The GB10 GPU is shared across GPU-consuming services. Only one GPU-heavy workload should run at a time:

Service	GPU Usage	Peak Memory	Typical Duration
Parabricks fq2bam	70-90% GPU	~40 GB	20-45 min
Parabricks DeepVariant	80-95% GPU	~60 GB	10-35 min
MolMIM NIM	Moderate	~8 GB	Always running
DiffDock NIM	Moderate	~8 GB	Always running
DCGM Exporter	Minimal	Minimal	Always running

## 8. Deploy Genomics Pipeline (Stage 1)

### 8.1 Parabricks Container Setup

#### BASH

```
# Pull Parabricks container for ARM64
docker pull nvcr.io/nvidia/clara/clara-parabricks:4.6.0-1

# Verify the image
docker images | grep parabricks
# Expected: clara-parabricks 4.6.0-1
```

### 8.2 BWA-MEM2 Alignment (fq2bam)

The [fq2bam](#) tool performs GPU-accelerated read alignment using BWA-MEM2 and produces a sorted, duplicate-marked BAM file.

#### BASH

```
docker run --rm --gpus all \
-v $(pwd)/genomics/data:/data \
nvcr.io/nvidia/clara/clara-parabricks:4.6.0-1 \
pbrun fq2bam \
--ref /data/reference/GRCh38.fa \
--in-fq /data/fastq/HG002_R1.fastq.gz /data/fastq/HG002_R2.fastq.gz \
--out-bam /data/bam/HG002.bam \
--num-gpus 1
```

#### Expected performance:

Metric	Value
Runtimes	20-45 minutes
GPU Utilization	70-90%
Peak GPU Memory	~40 GB
Output	Sorted, duplicate-marked BAM (~100 GB)

### 8.3 DeepVariant Variant Calling

#### BASH

```
docker run --rm --gpus all \
-v $(pwd)/genomics/data:/data \
nvcr.io/nvidia/clara/clara-parabricks:4.6.0-1 \
pbrun deepvariant \
--ref /data/reference/GRCh38.fa \
--in-bam /data/bam/HG002.bam \
--out-variants /data/vcf/HG002.vcf.gz \
--num-gpus 1
```

## Expected performance:

Metric	Value
Runtime	10-35 minutes
GPU Utilization	80-95%
Peak GPU Memory	~60 GB
Output	Compressed VCF (gzipped)

## 8.4 VCF Output Verification

### BASH

```
# Count total variants
zcat genomics/data/vcf/HG002.vcf.gz | grep -v '^#' | wc -l
# Expected: ~11,700,000 (11.7M variants)

# Count PASS variants with QUAL > 30
zcat genomics/data/vcf/HG002.vcf.gz | grep -v '^#' | \
awk '$7 == "PASS" && $6 > 30' | wc -l
# Expected: ~3,500,000 (3.5M)

# Count SNPs vs Indels
zcat genomics/data/vcf/HG002.vcf.gz | grep -v '^#' | \
awk '{if(length($4)==1 && length($5)==1) print "SNP"; else print "INDEL"}' | \
sort | uniq -c
# Expected: ~4,200,000 SNPs, ~1,000,000 indels
```

## VCF output summary:

Metric	Expected Value
Total variants	~11.7M
PASS variants (QUAL > 30)	~3.5M
SNPs	~4.2M
Indels	~1.0M
Coding region variants	~35,000

## 8.5 Genomics Portal (Port 5000)

After genomics processing, start the portal:

### BASH

```
docker compose up -d genomics-portal

# Verify
curl -s http://localhost:5000/health
# Expected: {"status": "healthy"}
```

Access the Genomics Portal at <http://<dgx-spark-ip>:5000> to browse VCF results.

## 8.6 Performance Benchmarks

Step	Wall Time	GPU Util	Peak Memory	Output Size
fq2bam (alignment)	20-45 min	70-90%	~40 GB	~100 GB BAM
DeepVariant (calling)	10-35 min	80-95%	~60 GB	~1 GB VCF.gz
<b>Total Stage 1</b>	<b>30-80 min</b>	—	—	—

## 9. Deploy RAG Chat Pipeline (Stage 2)

### 9.1 Milvus Vector Database Setup

#### BASH

```
# Start Milvus and its dependencies
docker compose up -d etcd minio milvus attu

# Wait for Milvus to be ready (30-60 seconds)
sleep 30

# Verify Milvus is running
curl -s http://localhost:19530/v1/health/ready
# Expected: {"status": "ok"}

# Verify Attu UI
curl -s -o /dev/null -w "%{http_code}" http://localhost:8000
# Expected: 200
```

### 9.2 Collection Schema

Create the `genomic_evidence` collection with 17 fields:

#### PYTHON

```
from pymilvus import connections, Collection, FieldSchema, CollectionSchema, DataType, utility

# Connect to Milvus
connections.connect(host="localhost", port=19530)

# Define schema with 17 fields
fields = [
    FieldSchema(name="id", dtype=DataType.INT64, is_primary=True, auto_id=True),
    FieldSchema(name="embedding", dtype=DataType.FLOAT_VECTOR, dim=384),
    FieldSchema(name="chrom", dtype=DataType.VARCHAR, max_length=10),
    FieldSchema(name="pos", dtype=DataType.INT64),
    FieldSchema(name="ref", dtype=DataType.VARCHAR, max_length=500),
    FieldSchema(name="alt", dtype=DataType.VARCHAR, max_length=500),
    FieldSchema(name="qual", dtype=DataType.FLOAT),
    FieldSchema(name="gene", dtype=DataType.VARCHAR, max_length=100),
    FieldSchema(name="consequence", dtype=DataType.VARCHAR, max_length=200),
    FieldSchema(name="impact", dtype=DataType.VARCHAR, max_length=20),
    FieldSchema(name="genotype", dtype=DataType.VARCHAR, max_length=10),
    FieldSchema(name="text_summary", dtype=DataType.VARCHAR, max_length=5000),
```

```

        FieldSchema(name="clinical_significance", dtype=DataType.VARCHAR, max_length=200),
        FieldSchema(name="rsid", dtype=DataType.VARCHAR, max_length=20),
        FieldSchema(name="disease_associations", dtype=DataType.VARCHAR, max_length=2000),
        FieldSchema(name="am_pathogenicity", dtype=DataType.FLOAT),
        FieldSchema(name="am_class", dtype=DataType.VARCHAR, max_length=20),
    ]

schema = CollectionSchema(fields, description="Genomic evidence for RAG")
collection = Collection("genomic_evidence", schema)

# Create IVF_FLAT index on embedding field
index_params = {
    "metric_type": "COSINE",
    "index_type": "IVF_FLAT",
    "params": {"nlist": 1024}
}
collection.create_index("embedding", index_params)

# Load collection into memory
collection.load()

print(f"Collection created: {collection.name}")
print(f"Schema fields: {len(fields)}")

```

#### Collection schema reference:

#	Field	Type	Details
1	id	INT64	Primary key, auto-generated
2	embedding	FLOAT_VECTOR	384 dimensions (BGE-small-en-v1.5)
3	chrom	VARCHAR(10)	Chromosome (chr1-22, chrX, chrY)
4	pos	INT64	Genomic position
5	ref	VARCHAR(500)	Reference allele
6	alt	VARCHAR(500)	Alternate allele
7	qual	FLOAT	Variant quality score
8	gene	VARCHAR(100)	Gene symbol
9	consequence	VARCHAR(200)	VEP functional consequence
10	impact	VARCHAR(20)	HIGH, MODERATE, LOW, MODIFIER
11	genotype	VARCHAR(10)	Sample genotype (e.g., 0/1, 1/1)
12	text_summary	VARCHAR(5000)	Natural-language variant summary
13	clinical_significance	VARCHAR(200)	ClinVar classification
14	rsid	VARCHAR(20)	dbSNP identifier
15	disease_associations	VARCHAR(2000)	Associated diseases/conditions
16	am_pathogenicity	FLOAT	AlphaMissense score (0.0-1.0)
17	am_class	VARCHAR(20)	pathogenic, ambiguous, or

## 9.3 Variant Annotation Pipeline

The annotation pipeline enriches VCF variants with data from three sources:

### BASH

```
# Run the annotation pipeline
python3 rag/annotation/clinvar.py \
--vcf genomics/data/vcf/HG002.vcf.gz \
--clinvar rag/data/clinvar/clinvar.vcf.gz \
--output rag/data/annotated_clinvar.tsv

python3 rag/annotation/alphamissense.py \
--vcf genomics/data/vcf/HG002.vcf.gz \
--am rag/data/alphamissense/AlphaMissense_hg38.tsv.gz \
--output rag/data/annotated_am.tsv

python3 rag/annotation/vep.py \
--vcf genomics/data/vcf/HG002.vcf.gz \
--output rag/data/annotated_vep.tsv
```

Expected annotation matches:

Source	Total Records	Patient Matches
ClinVar	4,100,000	~35,616
AlphaMissense	71,697,560	~6,831 (ClinVar-matched with predictions)
VEP	Per-variant	All coding variants

## 9.4 BGE Embedding and Indexing

### PYTHON

```
from sentence_transformers import SentenceTransformer
from pymilvus import connections, Collection

# Load embedding model
model = SentenceTransformer('BAAI/bge-small-en-v1.5') # 384 dimensions

# Connect to Milvus
connections.connect(host="localhost", port=19530)
collection = Collection("genomic_evidence")

# Example: embed and insert a variant
text = "chr9:35065263 G>A in VCP gene. ClinVar: Pathogenic. AlphaMissense: 0.87 (pathogenic). Consequence: missense_variant. Impact: MODERATE."
embedding = model.encode(text).tolist() # 384-dim vector

# Insert into Milvus
data = [
    {"embedding": embedding,
     "chrom": "chr9",
     "pos": 35065263,}
```

```

        "ref": "G",
        "alt": "A",
        "qual": 99.0,
        "gene": "VCP",
        "consequence": "missense_variant",
        "impact": "MODERATE",
        "genotype": "0/1",
        "text_summary": text,
        "clinical_significance": "Pathogenic",
        "rsid": "rs188935092",
        "disease_associations": "Inclusion body myopathy with Paget disease and frontotemporal dementia",
        "am_pathogenicity": 0.87,
        "am_class": "pathogenic"
    }]
}

collection.insert(data)
collection.flush()

```

#### Milvus index configuration:

Parameter	Value
Embedding Model	BGE-small-en-v1.5
Dimensions	384
Index Type	IVF_FLAT
Metric Type	COSINE
nlist	1024
nprobe (search)	16

## 9.5 Anthropic Claude Integration

### PYTHON

```

import anthropic

client = anthropic.Anthropic(api_key=os.environ["ANTHROPIC_API_KEY"])

def query_claude(question: str, context: str) -> str:
    """Send RAG query to Claude with retrieved genomic context."""
    response = client.messages.create(
        model="claude-sonnet-4-20250514",
        max_tokens=4096,
        temperature=0.3,
        messages=[{
            "role": "user",
            "content": f"""You are a genomics expert. Answer the question using the provided genomic
evidence.

Context:
{context}

Question: {question}"""
        }]
    )
    return response.content[0].text

```

## Claude configuration:

Parameter	Value
Model	claude-sonnet-4-20250514
Temperature	0.3
Max Tokens	4096

## 9.6 Knowledge Base

The platform includes a curated knowledge base of 201 genes across 13 therapeutic areas, with 171 genes (85%) classified as druggable.

Metric	Value
Total genes	201
Therapeutic areas	13
Druggable genes	171 (85%)

## 9.7 RAG API and Streamlit Chat

### BASH

```
# Start RAG API and Chat services
docker compose up -d rag-api streamlit-chat

# Verify RAG API
curl -s http://localhost:5001/health
# Expected: {"status": "healthy"}

# Verify Streamlit Chat
curl -s -o /dev/null -w "%{http_code}" http://localhost:8501
# Expected: 200
```

Access the Streamlit Chat at <http://<dgx-spark-ip>:8501> for conversational variant analysis.

# 10. Deploy Drug Discovery Pipeline (Stage 3)

## 10.1 BioNeMo NIM Services

### BASH

```
# Pull BioNeMo containers (requires NGC authentication)
docker pull nvcr.io/nvidia/clara/bionemo-molmim:1.0
docker pull nvcr.io/nvidia/clara/difffdock:1.0

# Start NIM services
docker compose up -d molmim difffdock

# Wait for models to load (may take 2-5 minutes)
sleep 120
```

```

# Verify MolMIM
curl -s http://localhost:8001/v1/health/ready
# Expected: {"status": "ready"}

# Verify DiffDock
curl -s http://localhost:8002/v1/health/ready
# Expected: {"status": "ready"}

```

## 10.2 10-Stage Pipeline Detail

Stage	Name	Input	Output	Key Operations
1	Initialize	Config + target gene	PipelineConfig	Validate parameters, create run ID
2	Normalize Target	Gene symbol	Normalized target	Map to UniProt, canonical name
3	Structure Discovery	UniProt ID	PDB structure list	Query RCSB PDB, score by resolution
4	Structure Preparation	PDB IDs	Prepared structures	Download PDB, extract binding sites
5	Molecule Generation	Seed SMILES + protein	Generated SMILES	MolMIM NIM (Port 8001)
6	Chemistry QC	SMILES list	Filtered SMILES	Lipinski, QED, TPSA checks
7	Conformer Generation	Filtered SMILES	3D conformers (SDF)	RDKit conformer embedding
8	Molecular Docking	Conformers + protein	Docking scores	DiffDock NIM (Port 8002)
9	Composite Ranking	All scores	Ranked candidates	Weighted composite formula
10	Reporting	Ranked candidates	PDF report	Visualizations, recommendations

## 10.3 Structure Retrieval and Scoring

### PYTHON

```

import requests

def search_pdb_structures(uniprot_id: str) -> list:
    """Search RCSB PDB for protein structures by UniProt ID."""
    url = "https://search.rcsb.org/rcsbsearch/v2/query"
    query = {
        "query": {
            "type": "terminal",
            "service": "text",
            "parameters": {
                "attribute": [
                    "rcsb_polymer_entity_container_identifiers.reference_sequence_identifiers.database_accession",
                    "operator": "exact_match",
                    "value": uniprot_id
                ]
            }
        }
    }
    response = requests.get(url, params=query)
    return response.json()

```

```

        }
    },
    "return_type": "entry"
}
response = requests.post(url, json=query)
return response.json().get("result_set", [])

```

## 10.4 Molecule Generation (MolMIM)

### PYTHON

```

import requests

def generate_molecules(seed_smiles: str, num_candidates: int = 100) -> list:
    """Generate molecule candidates using MolMIM NIM."""
    response = requests.post(
        "http://localhost:8001/generate",
        json={
            "smiles": seed_smiles,
            "num_molecules": num_candidates,
            "algorithm": "CMA-ES",
            "property_name": "QED",
            "min_similarity": 0.3,
            "particles": 30,
            "iterations": 10
        }
    )
    return response.json()["generated_molecules"]

```

## 10.5 Molecular Docking (DiffDock)

### PYTHON

```

def dock_molecule(protein_pdb: str, ligand_sdf: str) -> dict:
    """Score binding affinity using DiffDock NIM."""
    response = requests.post(
        "http://localhost:8002/molecular-docking/diffdock/generate",
        json={
            "protein": protein_pdb,
            "ligand": ligand_sdf,
            "num_poses": 10
        }
    )
    return response.json()

```

## 10.6 Drug-Likeness Scoring

Drug-likeness is assessed using three criteria:

### Lipinski Rule of Five:

Property	Threshold	Description
Molecular Weight	<= 500 Da	Size constraint
LogP	<= 5	Lipophilicity
H-Bond Donors (HBD)	<= 5	Polar surface groups

H-Bond Acceptors (HBA)

&lt;= 10

Polar surface groups

**Additional thresholds:**

Metric	Threshold	Interpretation
QED	> 0.67	Drug-like
TPSA	< 140 Angstrom squared	Good oral bioavailability

**PYTHON**

```
from rdkit import Chem
from rdkit.Chem import Descriptors, QED

def assess_drug_likeness(smiles: str) -> dict:
    """Evaluate drug-likeness using Lipinski, QED, and TPSA."""
    mol = Chem.MolFromSmiles(smiles)
    if mol is None:
        return {"valid": False}

    mw = Descriptors.MolWt(mol)
    logp = Descriptors.MolLogP(mol)
    hbd = Descriptors.NumHDonors(mol)
    hba = Descriptors.NumHAcceptors(mol)
    tpsa = Descriptors.TPSA(mol)
    qed_score = QED.qed(mol)

    lipinski_pass = (mw <= 500 and logp <= 5 and hbd <= 5 and hba <= 10)

    return {
        "valid": True,
        "mw": mw,
        "logp": logp,
        "hbd": hbd,
        "hba": hba,
        "tpsa": tpsa,
        "qed": qed_score,
        "lipinski_pass": lipinski_pass,
        "drug_like": qed_score > 0.67,
        "oral_bioavail": tpsa < 140
    }
```

## 10.7 Composite Ranking Formula

Candidates are ranked using a weighted composite score:

```
composite = 0.30 * generation_score + 0.40 * docking_score_normalized + 0.30 * qed_score
```

**Docking score normalization:****PYTHON**

```
def normalize_docking_score(dock_score: float) -> float:
    """Normalize docking score to [0, 1] range.
    More negative = better binding = higher normalized score."""
    return max(0.0, min(1.0, (10.0 + dock_score) / 20.0))
```

Raw Docking Score	Normalized Score	Interpretation
-10.0 kcal/mol	0.00	Excellent binding
-8.0 kcal/mol	0.10	Strong binding
-6.0 kcal/mol	0.20	Moderate binding
0.0 kcal/mol	0.50	Weak binding
+10.0 kcal/mol	1.00	No binding

**Note:** The normalization maps more negative (better) docking scores to lower normalized values. In the composite formula, the docking component rewards lower (better) scores.

#### Composite score weights:

Component	Weight	Source
Generation Score	30%	MolMIM similarity/property score
Docking Score (normalized)	40%	DiffDock binding affinity
QED Score	30%	RDKit quantitative drug-likeness

## 10.8 Discovery UI and Portal

### BASH

```
# Start Discovery services
docker compose up -d discovery-ui discovery-portal

# Verify Discovery UI
curl -s -o /dev/null -w "%{http_code}" http://localhost:8505
# Expected: 200

# Verify Discovery Portal
curl -s -o /dev/null -w "%{http_code}" http://localhost:8510
# Expected: 200
```

- **Discovery UI (Port 8505):** Interactive pipeline execution interface
- **Discovery Portal (Port 8510):** Results browser and reporting portal

## 10.9 PDF Report Generation

The final pipeline stage generates a PDF report containing:

- Target gene and variant summary
- PDB structure details with binding site analysis
- Top-ranked candidates with SMILES, scores, and 2D depictions
- Docking poses and binding affinity plots
- Lipinski and QED compliance table
- Composite score ranking

# 11. Nextflow Orchestration

## 11.1 DSL2 Pipeline Architecture

The HCLS AI Factory uses Nextflow DSL2 for pipeline orchestration. Each pipeline stage is defined as a separate process, with channels connecting inputs and outputs.

## 11.2 Pipeline Modes

Mode	Description	Stages Executed
full	Complete end-to-end pipeline	1 + 2 + 3 (all stages)
target	Start from target gene (skip genomics)	2 + 3
drug	Drug discovery only (pre-existing target)	3 only
demo	VCP demo with pre-loaded data	1 + 2 + 3 (demo subset)
genomics_only	Genomics pipeline only	1 only

## 11.3 Execution Profiles

Profile	Description	Use Case
standard	Local execution, default settings	Development
docker	Docker container execution	Standard deployment
singularity	Singularity container execution	HPC environments
dgx_spark	Optimized for DGX Spark hardware	Production on DGX Spark
slurm	SLURM workload manager	Multi-node clusters
test	Minimal test data, fast execution	CI/CD testing

## 11.4 Pipeline Launcher

### BASH

```
# Run with the pipeline launcher script
python3 scripts/run_pipeline.py \
  --mode full \
  --profile dgx_spark \
  --fastq genomics/data/fastq/ \
  --reference genomics/data/reference/GRCh38.fa

# Or run directly with Nextflow
nextflow run main.nf \
  -profile dgx_spark \
  --mode full \
  --fastq_dir genomics/data/fastq/ \
  --reference genomics/data/reference/GRCh38.fa \
  --outdir results/
```

## 11.5 Pipeline Configuration

### GROOVY

```
// nextflow.config
params {
    // Pipeline mode
    mode = 'full'

    // Input paths
    fastq_dir = 'genomics/data/fastq'
    reference = 'genomics/data/reference/GRCh38.fa'
    outdir = 'results'

    // Service endpoints
    milvus_host = 'localhost'
    milvus_port = 19530
    molmim_url = 'http://localhost:8001'
    diffdock_url = 'http://localhost:8002'

    // Drug discovery parameters
    num_candidates = 100
    min_qed = 0.67
    min_dock_score = -6.0
}

profiles {
    dgx_spark {
        docker.enabled = true
        docker.runOptions = '--gpus all'
        process {
            executor = 'local'
            memory = '120 GB'
            cpus = 128
        }
    }
}

test {
    params.mode = 'demo'
    process {
        memory = '16 GB'
        cpus = 4
    }
}
```

## 12. Service Startup and Health

### 12.1 start-services.sh Startup Order

Services should be started in dependency order:

### BASH

```
#!/bin/bash
# start-services.sh - Start all HCLS AI Factory services
```

```

set -e

echo "Starting infrastructure services..."
docker compose up -d etcd minio
sleep 10

echo "Starting Milvus..."
docker compose up -d milvus attu
sleep 30

echo "Starting BioNeMo NIM services..."
docker compose up -d molmim diffdock
sleep 120

echo "Starting application services..."
docker compose up -d genomics-portal rag-api streamlit-chat discovery-ui discovery-portal landing-page

echo "Starting monitoring..."
docker compose up -d prometheus grafana node-exporter dcgm-exporter

echo "All services started. Running health checks..."
sleep 10
bash scripts/validate_deployment.sh

```

## 12.2 Landing Page (Port 8080)

The landing page at <http://<dgx-spark-ip>:8080> provides a directory of all services with links and status indicators.

## 12.3 Health Check Endpoints

Service	Port	Health Endpoint	Expected Response
Genomics Portal	5000	/health	{"status": "healthy"}
RAG API	5001	/health	{"status": "healthy"}
Milvus	19530	/v1/health/ready	{"status": "ok"}
Attu	8000	/api/health	HTTP 200
Streamlit Chat	8501	/healthz	HTTP 200
MolMIM NIM	8001	/v1/health/ready	{"status": "ready"}
DiffDock NIM	8002	/v1/health/ready	{"status": "ready"}
Discovery UI	8505	/health	{"status": "healthy"}
Discovery Portal	8510	/health	{"status": "healthy"}
Grafana	3000	/api/health	{"status": "ok"}
Prometheus	9099	/-/healthy	HTTP 200
Node Exporter	9100	/metrics	Metrics text
DCGM Exporter	9400	/metrics	Metrics text

## 12.4 Verifying All Services

### BASH

```
#!/bin/bash
# validate_deployment.sh – Verify all services are running

declare -A SERVICES=(
    ["Landing Page"]="http://localhost:8080"
    ["Genomics Portal"]="http://localhost:5000/health"
    ["RAG API"]="http://localhost:5001/health"
    ["Milvus"]="http://localhost:19530/v1/health/ready"
    ["Attu"]="http://localhost:8000"
    ["Streamlit Chat"]="http://localhost:8501/healthz"
    ["MolMIM"]="http://localhost:8001/v1/health/ready"
    ["DiffDock"]="http://localhost:8002/v1/health/ready"
    ["Discovery UI"]="http://localhost:8505/health"
    ["Discovery Portal"]="http://localhost:8510/health"
    ["Grafana"]="http://localhost:3000/api/health"
    ["Prometheus"]="http://localhost:9099/-/healthy"
    ["Node Exporter"]="http://localhost:9100/metrics"
    ["DCGM Exporter"]="http://localhost:9400/metrics"
)

echo "==== HCLS AI Factory Health Check ==="
for service in "${!SERVICES[@]}"; do
    url="${SERVICES[$service]}"
    status=$(curl -s -o /dev/null -w "%{http_code}" "$url" 2>/dev/null || echo "ERR")
    if [ "$status" == "200" ]; then
        echo "[OK] $service ($url)"
    else
        echo "[FAIL] $service ($url) – HTTP $status"
    fi
done
```

## 13. Monitoring and Observability

### 13.1 Grafana Setup (Port 3000)

### BASH

```
# Start Grafana
docker compose up -d grafana

# Access at http://<dgx-spark-ip>:3000
# Default credentials: admin / changeme
```

Default Grafana credentials:

Parameter	Value
Username	admin
Password	changeme

## 13.2 Prometheus Configuration (Port 9099)

### YAML

```
# monitoring/prometheus/prometheus.yml
global:
  scrape_interval: 15s

scrape_configs:
  - job_name: 'node-exporter'
    static_configs:
      - targets: ['node-exporter:9100']

  - job_name: 'dcgm-exporter'
    static_configs:
      - targets: ['dcgm-exporter:9400']

  - job_name: 'rag-api'
    static_configs:
      - targets: ['rag-api:5001']
    metrics_path: /metrics

  - job_name: 'prometheus'
    static_configs:
      - targets: ['localhost:9090']
```

## 13.3 DCGM Exporter (Port 9400)

Key GPU metrics exposed by the DCGM Exporter:

Metric	Description
DCGM_FI_DEV_GPU_UTIL	GPU utilization percentage
DCGM_FI_DEV_FB_USED	GPU framebuffer memory used (MB)
DCGM_FI_DEV_FB_FREE	GPU framebuffer memory free (MB)
DCGM_FI_DEV_GPU_TEMP	GPU temperature (Celsius)
DCGM_FI_DEV_POWER_USAGE	Power consumption (Watts)
DCGM_FI_DEV_SM_CLOCK	Streaming multiprocessor clock (MHz)
DCGM_FI_DEV_MEM_CLOCK	Memory clock (MHz)

## 13.4 Node Exporter (Port 9100)

The Node Exporter provides host system metrics — CPU, memory, disk, and network utilization — critical for monitoring the DGX Spark ARM64 system.

## 13.5 Key Dashboard Panels

Recommended Grafana dashboard panels:

Panel	Data Source	Purpose
GPU Utilization	DCGM	Track fq2bam and DeepVariant GPU usage
GPU Memory	DCGM	Monitor peak memory during genomics
CPU Utilization	Node Exporter	ARM64 core usage across 144 cores
Memory Usage	Node Exporter	Unified 128 GB LPDDR5x utilization
Disk I/O	Node Exporter	NVMe throughput for FASTQ/BAM processing
Network I/O	Node Exporter	API call throughput
Container Status	Docker	Service health overview

## 13.6 Alert Configuration

### YAML

```
# Example alert rules for Prometheus
groups:
- name: hcls-alerts
  rules:
    - alert: GPUMemoryHigh
      expr: DCGM_FI_DEV_FB_USED / (DCGM_FI_DEV_FB_USED + DCGM_FI_DEV_FB_FREE) > 0.95
      for: 5m
      labels:
        severity: warning
      annotations:
        summary: "GPU memory usage above 95%"

    - alert: ServiceDown
      expr: up == 0
      for: 2m
      labels:
        severity: critical
      annotations:
        summary: "Service {{ $labels.job }} is down"
```

## 14. Security Configuration

### 14.1 API Key Management

#### BASH

```
# Store API keys in .env file (not committed to git)
echo ".env" >> .gitignore

# Set restrictive permissions
chmod 600 .env

# Verify .env is in .gitignore
grep -q '.env' .gitignore && echo "OK: .env is gitignored"
```

**Never commit API keys to version control.** Use environment variables exclusively:

Variable	Sensitivity	Storage
ANTHROPIC_API_KEY	High	.env file, chmod 600
NGC_API_KEY	High	.env file, chmod 600
GRAFANA_PASSWORD	Medium	.env file

## 14.2 Docker Network Isolation

Docker Compose creates an isolated bridge network. Only explicitly exposed ports are accessible from the host:

**BASH**

```
# Verify network isolation
docker network ls | grep hcls
docker network inspect hcls-ai-factory_default
```

## 14.3 Container Security

Best practices applied to the deployment:

- Run application containers as non-root users where possible
- Use read-only filesystem mounts for reference data
- Limit container capabilities with `--cap-drop ALL`
- Pin container image versions (no `latest` tags in production)

## 14.4 Data Access Controls

**BASH**

```
# Set appropriate permissions on data directories
chmod -R 750 genomics/data/
chmod -R 750 rag/data/
chmod -R 750 discovery/data/

# Ensure only the deployment user can access sensitive data
chown -R $(whoami):$(whoami) genomics/data/ rag/data/ discovery/data/
```

# 15. Data Management

## 15.1 Storage Layout

Directory	Contents	Size	Persistence
genomics/data/reference/	GRCh38 genome	3.1 GB	Permanent
genomics/data/fastq/	Input FASTQ files	~200 GB	Keep until processed
genomics/data/bam/	Alignment output	~100 GB	Delete after VCF

genomics/data/vcf/	Variant calls	~1 GB	Permanent
rag/data/clinvar/	ClinVar database	~1.2 GB	Permanent
rag/data/alphamissense/	AlphaMissense DB	~4 GB	Permanent
milvus_data (Docker volume)	Vector index	~2 GB	Permanent
discovery/data/	Structures, molecules	Variable	Per-run

## 15.2 Intermediate File Cleanup

BAM files are the largest intermediate output (~100 GB). Once the VCF has been verified, BAM files can be deleted to reclaim storage:

### BASH

```
# Verify VCF is complete before deleting BAM
zcat genomics/data/vcf/HG002.vcf.gz | grep -v '^#' | wc -l
# Confirm ~11.7M variants

# Delete intermediate BAM
rm -f genomics/data/bam/HG002.bam genomics/data/bam/HG002.bam.bai
echo "Reclaimed ~100 GB"
```

## 15.3 Milvus Data Persistence

Milvus data is stored in Docker volumes. To back up:

### BASH

```
# Stop Milvus for consistent backup
docker compose stop milvus

# Back up volumes
docker run --rm \
-v hcls-ai-factory_milvus_data:/data \
-v $(pwd)/backups:/backup \
alpine tar czf /backup/milvus_data_$(date +%Y%m%d).tar.gz /data

# Restart
docker compose start milvus
```

## 15.4 Backup Procedures

### BASH

```
# Full backup script
#!/bin/bash
BACKUP_DIR=~/backups/$(date +%Y%m%d)
mkdir -p $BACKUP_DIR

# Back up VCF results
cp -r genomics/data/vcf/ $BACKUP_DIR/vcf/

# Back up environment config (without secrets)
grep -v 'API_KEY' .env > $BACKUP_DIR/env_sanitized.txt
```

```

# Back up Milvus volumes
docker compose stop milvus
for vol in milvus_data etcd_data minio_data; do
  docker run --rm \
  -v hcls-ai-factory_${vol}:/data \
  -v $(pwd)/*$BACKUP_DIR:/backup \
  alpine tar czf /backup/${vol}.tar.gz /data
done
docker compose start milvus

echo "Backup complete: $BACKUP_DIR"

```

## 16. Performance Tuning

### 16.1 GPU Memory Management

The DGX Spark uses 128 GB unified LPDDR5x memory shared between CPU and GPU. Key considerations:

- Parabricks DeepVariant peaks at ~60 GB GPU memory — ensure other GPU services are idle during genomics processing
- MolMIM and DiffDock each require ~8 GB — they can co-exist during drug discovery
- Monitor with [nvidia-smi](#) and DCGM metrics during pipeline runs

#### BASH

```

# Monitor GPU memory in real-time
watch -n 1 nvidia-smi

# Check unified memory allocation
nvidia-smi --query-gpu=memory.used,memory.free,memory.total --format=csv

```

### 16.2 Milvus Index Tuning

Parameter	Default	Tuning Guidance
<code>nlist</code>	1024	Increase for larger collections (trade build time for search quality)
<code>nprobe</code>	16	Increase for higher recall (trade latency for accuracy)
<code>metric_type</code>	COSINE	Use COSINE for normalized BGE embeddings

#### PYTHON

```

# Search with tuned parameters
search_params = {
  "metric_type": "COSINE",
  "params": {"nprobe": 16}
}

```

```
results = collection.search(
    data=[query_embedding],
    anns_field="embedding",
    param=search_params,
    limit=10,
    output_fields=["gene", "clinical_significance", "text_summary"]
)
```

## 16.3 Docker Resource Limits

### YAML

```
# Example resource limits in docker-compose.yml
services:
  rag-api:
    deploy:
      resources:
        limits:
          memory: 16G
          cpus: '16'
        reservations:
          memory: 4G
          cpus: '4'
```

## 16.4 NVMe I/O Optimization

For FASTQ and BAM processing, I/O throughput is critical:

### BASH

```
# Check NVMe performance
fio --name=seqread --rw=read --bs=1M --size=1G --numjobs=4 --runtime=10 --group_reporting

# Ensure data directories are on NVMe
df -h genomics/data/
```

## 16.5 Pipeline Concurrency Settings

The Nextflow pipeline supports controlled concurrency:

### GROOVY

```
// nextflow.config - concurrency settings
process {
    maxForks = 4           // Maximum parallel processes
    maxRetries = 2          // Retry failed processes
    errorStrategy = 'retry'
}

executor {
    queueSize = 8          // Maximum queued tasks
    pollInterval = '5 sec'
}
```

# 17. Troubleshooting Guide

## 17.1 Service Not Starting

### BASH

```
# Check service logs
docker compose logs <service-name> --tail 50

# Check if port is already in use
ss -tlnp | grep <port>

# Restart a specific service
docker compose restart <service-name>
```

## 17.2 GPU Out of Memory

### BASH

```
# Check current GPU memory usage
nvidia-smi

# Kill any orphaned GPU processes
sudo fuser -v /dev/nvidia*

# Reduce Parabricks memory by limiting GPU threads
# Add --gpu-mem-limit flag if available

# Ensure NIM services are stopped during genomics
docker compose stop molmim diffdock
```

## 17.3 Milvus Connection Issues

### BASH

```
# Verify Milvus dependencies are running
docker compose ps etcd minio milvus

# Check Milvus logs for errors
docker compose logs milvus --tail 100

# Test connectivity
curl -s http://localhost:19530/v1/health/ready

# Reset Milvus if corrupted
docker compose down milvus etcd minio
docker volume rm hcls-ai-factory_milvus_data hcls-ai-factory_etcd_data hcls-ai-factory_minio_data
docker compose up -d etcd minio milvus
```

## 17.4 BioNeMo NIM Not Ready

### BASH

```
# NIM services may take 2-5 minutes to load models
# Check logs for model loading progress
docker compose logs molmim --tail 50
docker compose logs diffdock --tail 50

# Verify GPU is available for NIM
nvidia-smi | grep -i "molmim|diffdock"

# Restart if stuck
docker compose restart molmim diffdock
```

## 17.5 Parabricks Failures

Error	Cause	Resolution
CUDA out of memory	Insufficient GPU memory	Stop other GPU services first
Reference index not found	Missing .fai file	Run <a href="#">samtools faidx GRCh38.fa</a>
Input file not found	Wrong FASTQ path	Check volume mount paths
Unsupported GPU	Driver mismatch	Update NVIDIA driver

## 17.6 Claude API Errors

Error	Cause	Resolution
401 Unauthorized	Invalid API key	Verify <a href="#">ANTHROPIC_API_KEY</a> in .env
429 Rate Limited	Too many requests	Implement exponential backoff
500 Server Error	Anthropic service issue	Retry after 30 seconds
Connection refused	No internet	Check network connectivity

## 17.7 Docker Issues

### BASH

```
# Docker daemon not running
sudo systemctl start docker
sudo systemctl enable docker

# Disk space full
docker system prune -a --volumes
df -h /var/lib/docker

# Permission denied
sudo usermod -aG docker $USER
newgrp docker
```

## 17.8 Common Error Messages Table

Error Message	Service	Resolution
Connection refused on port 19530	Milvus	Start etcd + MinIO first, then Milvus
NVIDIA driver not found	Docker	Install NVIDIA Container Toolkit
Model not loaded	MolMIM/DiffDock	Wait 2-5 minutes for model loading
Collection not found	Milvus	Run schema creation script (Section 9.2)
API key not set	RAG API	Set <code>ANTHROPIC_API_KEY</code> in <code>.env</code>
Out of disk space	Parabicks	Clean BAM intermediates, expand storage
Permission denied: /data	Any	Check volume mount permissions

## 18. VCP/FTD Demo Walkthrough

### 18.1 Demo Overview

The VCP (Valosin-Containing Protein) / FTD (Frontotemporal Dementia) demo showcases the full three-stage pipeline using a known pathogenic variant:

Parameter	Value
Variant	rs188935092
Location	chr9:35065263 G>A
Gene	VCP
ClinVar Classification	Pathogenic
AlphaMissense Score	0.87 (pathogenic, threshold >0.564)
Disease	Inclusion body myopathy with Paget disease and FTD
Seed Molecule	CB-5083 (VCP/p97 inhibitor)
PDB Structures	8OOI, 9DIL, 7K56, 5FTK
Binding Domain	D2 ATPase domain, ~450 cubic angstroms
Druggability Score	0.92

### 18.2 Pre-Demo Setup

#### BASH

```
# Ensure all services are running
bash scripts/validate_deployment.sh

# Verify Milvus has the VCP variant loaded
python3 -c "
from pymilvus import connections, Collection
connections.connect(host='localhost', port=19530)
col = Collection('genomic_evidence')
```

```

col.load()
results = col.query('gene == \'VCP\'', output_fields=['rsid', 'clinical_significance', 'am_pathogenicity'])
print(f'VCP variants found: {len(results)}')
for r in results[:3]:
    print(r)
"
```

## 18.3 Running the Demo

### BASH

```

# Run the demo pipeline mode
python3 scripts/run_pipeline.py --mode demo

# Or via Nextflow
nextflow run main.nf -profile dgx_spark --mode demo
```

### Step-by-step execution:

- 1. Stage 1 (Genomics):** Process demo FASTQ subset through Parabricks fq2bam and DeepVariant
- 2. Stage 2 (RAG):** Annotate VCP variant with ClinVar (Pathogenic) and AlphaMissense (0.87), embed into Milvus, query Claude for clinical interpretation
- 3. Stage 3 (Drug Discovery):** Retrieve PDB structures (8OOI, 9DIL, 7K56, 5FTK), generate molecules from CB-5083 seed via MolMIM, dock with DiffDock, rank by composite score

## 18.4 Expected Results

Metric	Expected Value
Candidates generated	100
Pass Lipinski Rule of Five	87
QED > 0.67 (drug-like)	72
Top docking scores	-8.2 to -11.4 kcal/mol
Composite score range	0.68 - 0.89

### Top candidate characteristics:

Property	Range
Molecular Weight	300 - 500 Da
LogP	1.5 - 4.5
QED	0.67 - 0.92
TPSA	40 - 130 squared angstroms
Docking Score	-8.2 to -11.4 kcal/mol
Composite Score	0.68 - 0.89

# 19. Scaling Beyond DGX Spark

## 19.1 Phase 1 to Phase 3 Roadmap

Phase	Hardware	Scale	Use Case
Phase 1	DGX Spark	Single workstation	Development, demos, single-patient analysis
Phase 2	DGX B200	Single server, multi-GPU	Production cohort analysis
Phase 3	DGX SuperPOD	Multi-node cluster	Population-scale genomics

## 19.2 Kubernetes Migration Path

For Phase 2 and beyond, migrate from Docker Compose to Kubernetes:

- Replace `docker-compose.yml` with Helm charts
- Use NVIDIA GPU Operator for GPU scheduling
- Deploy Milvus Cluster mode (distributed) instead of standalone
- Use persistent volume claims (PVCs) for data storage
- Implement horizontal pod autoscaling for RAG API

## 19.3 Multi-GPU Considerations

- Parabricks supports `--num-gpus` for multi-GPU parallelism
- MolMIM and DiffDock can be replicated across GPUs
- Milvus supports distributed deployment with multiple query nodes

## 19.4 NVIDIA FLARE for Federated Learning

For multi-institutional deployments, NVIDIA FLARE enables federated learning across DGX Spark nodes without sharing raw patient data.

# 20. Appendix A: Complete Configuration Reference

## 20.1 All Environment Variables

Variable	Default	Description
<code>ANTHROPIC_API_KEY</code>	(required)	Anthropic API key for Claude
<code>NGC_API_KEY</code>	(required)	NVIDIA NGC API key
<code>REFERENCE_GENOME</code>	<code>/data/reference/GRCh38.fa</code>	Path to reference genome
<code>MILVUS_HOST</code>	<code>localhost</code>	Milvus server hostname
<code>MILVUS_PORT</code>	<code>19530</code>	Milvus server port
<code>MOLMIM_URL</code>	<code>http://localhost:8001</code>	MolMIM NIM endpoint

DIFFDOCK_URL	<a href="http://localhost:8002">http://localhost:8002</a>	DiffDock NIM endpoint
CLAUDE_MODEL	claude-sonnet-4-20250514	Claude model identifier
CLAUDE_TEMPERATURE	0.3	Claude sampling temperature
PIPELINE_MODE	full	Pipeline execution mode
NUM_CANDIDATES	100	Number of molecules to generate
MIN_QED	0.67	Minimum QED threshold
MIN_DOCK_SCORE	-6.0	Minimum docking score (kcal/mol)
GRAFANA_USER	admin	Grafana admin username
GRAFANA_PASSWORD	changeme	Grafana admin password

## 20.2 AlphaMissense Thresholds

Classification	Score Range
Pathogenic	> 0.564
Ambiguous	0.34 - 0.564
Benign	< 0.34

## 20.3 Scoring Weights

Component	Weight
Generation Score	0.30 (30%)
Docking Score (normalized)	0.40 (40%)
QED Score	0.30 (30%)

## 20.4 Drug-Likeness Thresholds

Property	Threshold	Rule
Molecular Weight	<= 500 Da	Lipinski
LogP	<= 5	Lipinski
H-Bond Donors	<= 5	Lipinski
H-Bond Acceptors	<= 10	Lipinski
QED	> 0.67	Drug-likeness
TPSA	< 140 squared angstroms	Oral bioavailability

## 20.5 Docking Score Interpretation

Score (kcal/mol)	Binding Affinity	Assessment
< -10.0	Excellent	Strong candidate
-8.0 to -10.0	Strong	Viable candidate
-6.0 to -8.0	Moderate	Marginal candidate

> -6.0

Weak

Poor candidate

#### Normalization formula:

```
normalized = max(0, min(1, (10 + dock_score) / 20))
```

## 21. Appendix B: API Reference

### 21.1 MolMIM API (Port 8001)

#### Generate Molecules:

##### JSON

```
// POST http://localhost:8001/generate
// Request:
{
  "smiles": "CC1=CC=C(C=C1)C(=O)NC2=CC=CC=C2",
  "num_molecules": 100,
  "algorithm": "CMA-ES",
  "property_name": "QED",
  "min_similarity": 0.3,
  "particles": 30,
  "iterations": 10
}

// Response:
{
  "generated_molecules": [
    {
      "smiles": "CC1=CC=C(C=C1)C(=O)NC2=CC=C(F)C=C2",
      "score": 0.85,
      "similarity": 0.78
    }
  ]
}
```

#### Health Check:

```
GET http://localhost:8001/v1/health/ready
Response: {"status": "ready"}
```

### 21.2 DiffDock API (Port 8002)

#### Molecular Docking:

##### JSON

```
// POST http://localhost:8002/molecular-docking/difffdock/generate
// Request:
{
  "protein": "<PDB file content>",


```

```

    "ligand": "<SDF file content>",
    "num_poses": 10
}

// Response:
{
  "poses": [
    {
      "pose_id": 0,
      "confidence": 0.95,
      "score": -9.7,
      "ligand_sdf": "<docked SDF content>"
    }
  ]
}

```

#### Health Check:

```

GET http://localhost:8002/v1/health/ready
Response: {"status": "ready"}

```

## 21.3 RAG API Endpoints (Port 5001)

Method	Endpoint	Description
GET	<a href="#">/health</a>	Service health check
POST	<a href="#">/query</a>	RAG query with context retrieval
POST	<a href="#">/search</a>	Vector similarity search
GET	<a href="#">/collections</a>	List Milvus collections
GET	<a href="#">/stats</a>	Collection statistics

#### RAG Query Example:

##### JSON

```

// POST http://localhost:5001/query
// Request:
{
  "question": "What pathogenic variants are found in the VCP gene?",
  "top_k": 10,
  "filters": {
    "gene": "VCP",
    "impact": "HIGH"
  }
}

// Response:
{
  "answer": "The VCP gene contains the variant rs188935092...",
  "sources": [
    {
      "gene": "VCP",
      "rsid": "rs188935092",
      "clinical_significance": "Pathogenic",
      "am_pathogenicity": 0.87,
      "similarity_score": 0.94
    }
  ]
}

```

```
],
  "model": "claude-sonnet-4-20250514",
  "tokens_used": 1847
}
```

## 21.4 Health Check Endpoints Summary

Service	Endpoint	Method
Genomics Portal	/health	GET
RAG API	/health	GET
Milvus	/v1/health/ready	GET
Attu	/api/health	GET
Streamlit Chat	/healthz	GET
MolMIM	/v1/health/ready	GET
DiffDock	/v1/health/ready	GET
Discovery UI	/health	GET
Discovery Portal	/health	GET
Grafana	/api/health	GET
Prometheus	/-/healthy	GET
Node Exporter	/metrics	GET
DCGM Exporter	/metrics	GET

## 22. Appendix C: Schema Definitions

### 22.1 Milvus Collection Schema

Collection: `genomic\_evidence`

#	Field	Data Type	Constraints	Description
1	id	INT64	Primary Key, Auto ID	Unique record identifier
2	embedding	FLOAT_VECTOR	dim=384	BGE-small-en-v1.5 embedding
3	chrom	VARCHAR	max_length=10	Chromosome (chr1-22, chrX, chrY)
4	pos	INT64	—	Genomic position (1-based)
5	ref	VARCHAR	max_length=500	Reference allele
6	alt	VARCHAR	max_length=500	Alternate allele
7	qual	FLOAT	—	Variant quality score
8	gene	VARCHAR	max_length=100	HGNC gene symbol
9	consequence	VARCHAR	max_length=200	VEP consequence term
10	impact	VARCHAR	max_length=20	HIGH/MODERATE/LOW/MODIFIER
11	genotype	VARCHAR	max_length=10	Sample genotype (0/1, 1/1)
12	text_summary	VARCHAR	max_length=5000	Natural-language summary

13	<code>clinical_significance</code>	VARCHAR	<code>max_length=200</code>	ClinVar classification
14	<code>rsid</code>	VARCHAR	<code>max_length=20</code>	dbSNP RS identifier
15	<code>disease_associations</code>	VARCHAR	<code>max_length=2000</code>	Associated diseases
16	<code>am_pathogenicity</code>	FLOAT	<code>0.0-1.0</code>	AlphaMissense pathogenicity
17	<code>am_class</code>	VARCHAR	<code>max_length=20</code>	pathogenic/ambiguous/benign

#### Index configuration:

Parameter	Value
Index Type	IVF_FLAT
Metric Type	COSINE
nlist	1024
nprobe (search)	16

## 22.2 Pydantic Data Models

### PYTHON

```
from pydantic import BaseModel, Field
from typing import List, Optional
from enum import Enum

class TargetHypothesis(BaseModel):
    """Genomic target identified from variant analysis."""
    gene: str
    variant_id: str
    rsid: Optional[str]
    clinical_significance: str
    am_pathogenicity: Optional[float]
    am_class: Optional[str]
    therapeutic_area: str
    druggability_score: float
    rationale: str

class StructureInfo(BaseModel):
    """PDB structure information for a target protein."""
    pdb_id: str
    resolution: float
    method: str
    chain: str
    binding_site_volume: Optional[float]

class StructureManifest(BaseModel):
    """Collection of structures for a target."""
    target_gene: str
    uniprot_id: str
    structures: List[StructureInfo]
    selected_structure: str

class MoleculeProperties(BaseModel):
    """Chemical properties of a generated molecule."""
    molecular_weight: float
    logp: float
    hbd: int
    hba: int
```

```

        tpsa: float
        qed: float
        lipinski_pass: bool

class GeneratedMolecule(BaseModel):
    """Molecule generated by MolMIM."""
    smiles: str
    generation_score: float
    similarity_to_seed: float
    properties: MoleculeProperties

class DockingResult(BaseModel):
    """Molecular docking result from DiffDock."""
    smiles: str
    dock_score: float # kcal/mol (negative = better)
    confidence: float
    pose_sdf: str

class RankedCandidate(BaseModel):
    """Final ranked drug candidate with composite score."""
    rank: int
    smiles: str
    generation_score: float
    dock_score: float
    dock_score_normalized: float
    qed: float
    composite_score: float # 0.3*gen + 0.4*dock + 0.3*qed
    lipinski_pass: bool
    properties: MoleculeProperties

class PipelineConfig(BaseModel):
    """Configuration for a pipeline run."""
    mode: str = "full"
    target_gene: Optional[str]
    seed_smiles: Optional[str]
    num_candidates: int = 100
    min_qed: float = 0.67
    min_dock_score: float = -6.0
... (16 more lines)

```

## 23. Appendix D: Docker Image Reference

### 23.1 All Container Images

Service	Image	Tag	Architecture
Parabricks	nvcr.io/nvidia/clara/clara-parabricks	4.6.0-1	ARM64 (aarch64)
Milvus	milvusdb/milvus	v2.4-latest	ARM64
MolMIM	nvcr.io/nvidia/clara/bionemo-molmim	1.0	ARM64
DiffDock	nvcr.io/nvidia/clara/difffdock	1.0	ARM64

Grafana	grafana/grafana	10.2.2	ARM64
Prometheus	prom/prometheus	v2.48.0	ARM64
Node Exporter	prom/node-exporter	latest	ARM64
DCGM Exporter	nvcr.io/nvidia/k8s/dcgm-exporter	latest	ARM64
etcd	quay.io/coreos/etcd	v3.5.5	ARM64
MinIO	minio/minio	latest	ARM64
Attu	zilliz/attu	latest	ARM64

## 23.2 ARM64 Compatibility Notes

The DGX Spark uses an ARM64 (aarch64) processor. All container images must be ARM64-compatible:

- NVIDIA NGC images for Parabricks, BioNeMo, and DCGM include ARM64 variants
- Community images (Grafana, Prometheus, MinIO, etcd) provide multi-arch manifests
- Custom application images must be built with `--platform linux/arm64`
- If building locally, ensure the base image supports ARM64

### BASH

```
# Verify image architecture
docker inspect --format='{{.Architecture}}' <image-name>
# Expected: arm64

# Build for ARM64 explicitly
docker build --platform linux/arm64 -t my-service:latest ./my-service/
```

## 24. Appendix E: Validation Checklists

### 24.1 Pre-Deployment Checklist

#	Item	Command / Check	Expected
1	DGX Spark hardware	<code>uname -m</code>	aarch64
2	GPU detected	<code>nvidia-smi</code>	GB10 GPU listed
3	Docker installed	<code>docker --version</code>	24.0+
4	Docker Compose V2	<code>docker compose version</code>	v2.x
5	NVIDIA runtime	<code>`docker info \  grep nvidia`</code>	nvidia listed
6	Python version	<code>python3 --version</code>	3.10+
7	Disk space	<code>df -h /</code>	>= 320 GB free
8	Reference genome	<code>ls genomics/data/reference/GRCh38.fa</code>	File exists, ~3.1 GB
9	ClinVar data	<code>ls rag/data/clinvar/clinvar.vcf.gz</code>	File exists,

			~1.2 GB
10	AlphaMissense data	<code>ls rag/data/alphamissense/AlphaMissense_hg38.tsv.gz</code>	File exists, ~4 GB
11	API keys configured	<code>grep ANTHROPIC_API_KEY .env</code>	Key set (not empty)
12	NGC key configured	<code>grep NGC_API_KEY .env</code>	Key set (not empty)
13	.env permissions	<code>stat -c %a .env</code>	600
14	.env in .gitignore	<code>grep .env .gitignore</code>	Present

## 24.2 Post-Deployment Checklist

#	Item	Command / Check	Expected
1	All containers running	<code>docker compose ps</code>	14+ services "Up"
2	Landing Page	<code>curl http://localhost:8080</code>	HTTP 200
3	Genomics Portal	<code>curl http://localhost:5000/health</code>	{"status": "healthy"}
4	RAG API	<code>curl http://localhost:5001/health</code>	{"status": "healthy"}
5	Milvus ready	<code>curl http://localhost:19530/v1/health/ready</code>	{"status": "ok"}
6	Attu UI	<code>curl -o /dev/null -w "%{http_code}" http://localhost:8000</code>	200
7	Streamlit Chat	<code>curl -o /dev/null -w "%{http_code}" http://localhost:8501</code>	200
8	MolMIM ready	<code>curl http://localhost:8001/v1/health/ready</code>	{"status": "ready"}
9	DiffDock ready	<code>curl http://localhost:8002/v1/health/ready</code>	{"status": "ready"}
10	Discovery UI	<code>curl http://localhost:8505/health</code>	{"status": "healthy"}
11	Discovery Portal	<code>curl http://localhost:8510/health</code>	{"status": "healthy"}
12	Grafana	<code>curl http://localhost:3000/api/health</code>	{"status": "ok"}
13	Prometheus	<code>curl http://localhost:9099/-/healthy</code>	HTTP 200
14	DCGM metrics	<code>curl http://localhost:9400/metrics</code>	Metrics text
15	Milvus collection	Python: <code>Collection("genomic_evidence").num_entities</code>	> 0

## 24.3 Demo Readiness Checklist

#	Item	Check	Expected
1	All services healthy	Run <code>validate_deployment.sh</code>	All [OK]
2	VCP variant in Milvus	Query gene="VCP"	rs188935092 found
3	ClinVar annotation	VCP classification	Pathogenic
4	AlphaMissense score	VCP am_pathogenicity	0.87
5	PDB structures accessible	Query RCSB for VCP	800I, 9DIL, 7K56, 5FTK
6	MolMIM generates	Test generation from CB-	Molecules returned

		5083	
7	DiffDock docks	Test docking against VCP structure	Scores returned
8	Claude responds	Test RAG query about VCP	Coherent response
9	Grafana dashboards	Login at port 3000	Dashboards visible
10	GPU metrics flowing	Check DCGM in Grafana	GPU util, memory shown

## 25. Appendix F: Glossary

### 25.1 Genomics Terms

Term	Definition
<b>FASTQ</b>	Text-based format for storing nucleotide sequences and quality scores
<b>BAM</b>	Binary Alignment Map — compressed format for aligned sequencing reads
<b>VCF</b>	Variant Call Format — standard format for genomic variants
<b>SNP</b>	Single Nucleotide Polymorphism — single base-pair variant
<b>Indel</b>	Insertion or deletion of nucleotides in the genome
<b>WGS</b>	Whole Genome Sequencing — sequencing of entire genome
<b>GRCh38</b>	Genome Reference Consortium Human Build 38 — current reference genome
<b>GIAB</b>	Genome in a Bottle — NIST benchmark samples (e.g., HG002)
<b>ClinVar</b>	NCBI database of clinically relevant genomic variants
<b>VEP</b>	Variant Effect Predictor — functional annotation tool
<b>AlphaMissense</b>	DeepMind model predicting missense variant pathogenicity
<b>Paired-end</b>	Sequencing both ends of a DNA fragment for improved alignment
<b>Coverage (30x)</b>	Average number of reads covering each position in the genome

### 25.2 ML/AI Terms

Term	Definition
<b>RAG</b>	Retrieval-Augmented Generation — combining search with LLM generation
<b>Embedding</b>	Dense vector representation of text or data
<b>BGE</b>	BAAI General Embedding — sentence transformer model family
<b>IVF_FLAT</b>	Inverted File Index — approximate nearest neighbor search method

<b>COSINE</b>	Cosine similarity — metric for comparing vector directions
<b>NIM</b>	NVIDIA Inference Microservice — containerized model serving
<b>LLM</b>	Large Language Model — e.g., Claude
<b>Vector Database</b>	Database optimized for similarity search on dense vectors
<b>nlist</b>	Number of clusters in IVF index (build-time parameter)
<b>nprobe</b>	Number of clusters to search at query time (recall vs. latency)

## 25.3 Drug Discovery Terms

Term	Definition
<b>SMILES</b>	Simplified Molecular Input Line Entry System — text notation for molecules
<b>PDB</b>	Protein Data Bank — repository of 3D protein structures
<b>Molecular Docking</b>	Computational prediction of ligand-protein binding pose and affinity
<b>QED</b>	Quantitative Estimate of Drug-likeness — composite drug-likeness score (0-1)
<b>Lipinski Rule of Five</b>	Empirical rules predicting oral bioavailability
<b>TPSA</b>	Topological Polar Surface Area — predictor of membrane permeability
<b>LogP</b>	Partition coefficient — measure of lipophilicity
<b>HBD / HBA</b>	Hydrogen Bond Donors / Acceptors
<b>Conformer</b>	3D spatial arrangement of a molecule's atoms
<b>Binding Affinity</b>	Strength of interaction between a drug molecule and its target protein
<b>kcal/mol</b>	Kilocalories per mole — unit for binding energy (more negative = stronger)
<b>MoLMIM</b>	Molecule generation model from NVIDIA BioNeMo
<b>DiffDock</b>	Diffusion-based molecular docking model
<b>Druggability</b>	Assessment of whether a protein target can be modulated by a small molecule
<b>CB-5083</b>	VCP/p97 inhibitor used as seed molecule in the VCP demo
<b>RDKit</b>	Open-source cheminformatics toolkit for molecular analysis

*This deployment guide is maintained as part of the HCLS AI Factory open-source project. For updates, issues, and contributions, visit the project repository on GitHub.*