

Executive Bullets

HCLS AI Factory

From Patient DNA to Novel Drug Candidates in < 5 Hours

NVIDIA DGX Spark | Parabricks | BioNeMo | Milvus | Claude

02/2026 | Version 1.0 | Apache 2.0 License | Author: Adam Jones

What It Is

The HCLS AI Factory transforms patient DNA into ranked novel drug candidates in under 5 hours on a single NVIDIA DGX Spark (\$3,999). Three GPU-accelerated stages run end-to-end with no manual intervention.

The Problem

- CPU-based genomics: 12-36 hours for a single 30x WGS sample
- Variant annotation fragmented across disconnected databases
- Variant-to-drug-lead gap: months of manual work
- No integrated platform connects genomics → reasoning → drug discovery

The Solution — Three Stages

Stage 1: GPU-Accelerated Genomics (120-240 min)

- NVIDIA Parabricks 4.6 — 10-20x faster than CPU
- BWA-MEM2 alignment: 20-45 min | DeepVariant: 10-35 min, >99% accuracy
- Input: ~200 GB FASTQ | Output: VCF (~11.7M variants)

Stage 2: RAG-Grounded Target Identification (Interactive)

- ClinVar (4.1M) + AlphaMissense (71M) + VEP annotation
- 3.5M variant embeddings in Milvus | BGE-small-en-v1.5 (384-dim)
- Claude RAG reasoning | 201 genes, 13 therapeutic areas, 171 druggable (85%)

Stage 3: AI-Driven Drug Discovery (8-16 min)

- BioNeMo MolMIM (generation) + DiffDock (docking) + RDKit (scoring)
- Composite: 30% generation + 40% docking + 30% QED
- Output: 100 ranked novel drug candidates + PDF report

Key Numbers

| Metric | Value |
|-----------------------|---------------------------------|
| Total Pipeline Time | < 5 hours |
| Input Data | ~200 GB FASTQ (30x WGS) |
| Variants Called | ~11.7 million |
| High-Quality Variants | ~3.5 million (QUAL>30) |
| Knowledge Base | 201 genes, 13 therapeutic areas |
| Druggable Targets | 171 (85%) |
| Drug Candidates | 100 (ranked by composite score) |
| Hardware Cost | \$3,999 (DGX Spark) |

VCP/FTD Demo Highlights

- Target: VCP gene — FTD, ALS, IBMPFD
- Variant: rs188935092 — ClinVar Pathogenic, AlphaMissense 0.87
- Seed: CB-5083 (Phase I VCP inhibitor)
- Top candidate: +39% composite improvement over seed
- Docking: -11.4 kcal/mol (vs. -8.1) | QED: 0.81 (vs. 0.62)
- All top 10 pass Lipinski's Rule of Five

Technology Stack

| Layer | Technology |
|----------------|--|
| Hardware | NVIDIA DGX Spark (GB10, 128 GB unified, \$3,999) |
| Genomics | Parabricks 4.6, DeepVariant (>99% accuracy) |
| Annotation | ClinVar, AlphaMissense, Ensembl VEP |
| Vector DB | Milvus 2.4, BGE-small-en-v1.5, IVF_FLAT |
| LLM | Anthropic Claude (RAG-grounded reasoning) |
| Drug Discovery | BioNeMo MolMIM, DiffDock, RDKit |
| Orchestration | Nextflow DSL2 (5 modes) |
| Monitoring | Grafana, Prometheus, DCGM Exporter |
| License | Apache 2.0 (fully open) |

Deployment Roadmap

| Phase | Hardware | Scale | Cost |
|------------------|--------------|----------------------------|-------------|
| 1 — Proof Build | DGX Spark | 1 patient, Docker Compose | \$3,999 |
| 2 — Departmental | DGX B200 | Multiple concurrent, K8s | \$500K-\$1M |
| 3 — Enterprise | DGX SuperPOD | Thousands, FLARE federated | \$7M-\$60M+ |

Cross-Modal Integration

- Imaging → Genomics: Lung-RADS 4B+ triggers tumor profiling
- Genomics → Drug Discovery: Pathogenic variants trigger molecule generation
- NVIDIA FLARE: Federated learning (data stays local)

Competitive Differentiation

- Only platform: genomics-to-drug-candidates on a single desktop GPU
- End-to-end: No manual handoffs between stages
- < 5 hours total (vs. weeks/months traditional)
- \$3,999 proof build (vs. \$100K+ CPU infrastructure)
- Open-source: Apache 2.0, reproducible, auditable
- Scalable: Same pipelines from DGX Spark to SuperPOD

HCLS AI Factory — Apache 2.0 / Author: Adam Jones / February 2026