



**Learning Guide: Advanced**

# HCLS AI Factory

## Graduate / Professional Level

Deep technical analysis of the HCLS AI Factory architecture, from BWA-MEM2 seed-and-extend algorithms through diffusion-based molecular docking, with emphasis on algorithmic design decisions, scaling bottlenecks, and clinical translation barriers.

*NVIDIA DGX Spark / Parabricks / BioNeMo / Milvus / Claude*

02/2026 | Version 1.0 | Apache 2.0 License

Author: Adam Jones

# Table of Contents

1. Computational Genomics — From FASTQ to VCF
  2. Variant Annotation — Multi-Database Integration
  3. Vector Database Architecture — Milvus and RAG
  4. Drug Discovery Pipeline — Deep Dive
  5. Nextflow DSL2 Pipeline Architecture
  6. Clinical Translation and Limitations
  7. Scaling Analysis
  8. Advanced Topics and Extensions
- Review Questions

# Chapter 1: Computational Genomics

## — From FASTQ to VCF

### 1.1 Sequencing Data Characteristics

The HCLS AI Factory processes Illumina short-read data: 2×250 bp paired-end reads from 30× whole-genome sequencing of HG002 (NA24385), a GIAB Ashkenazi Jewish reference standard. The FASTQ files total approximately 200 GB and contain ~800 million read pairs.

#### Why HG002?

The Genome in a Bottle (GIAB) Consortium provides extensively validated truth sets for HG002, enabling rigorous benchmarking. The high-confidence regions cover >95% of the GRCh38 reference, with variant calls validated by multiple orthogonal technologies (PacBio HiFi, Oxford Nanopore, Hi-C, optical mapping).

### 1.2 GPU-Accelerated Alignment: BWA-MEM2 on Parabicks

NVIDIA Parabicks 4.6.0-1 (container: nvcr.io/nvidia/clara/clara-parabicks:4.6.0-1) provides a GPU-accelerated implementation of BWA-MEM2.

#### Algorithm Overview

BWA-MEM2 uses a seed-and-extend approach:

- 1. Seeding:** Extract fixed-length k-mers from the query read and look them up in the FM-index of the reference genome
- 2. Chaining:** Group collinear seeds into chains representing candidate alignment locations
- 3. Extension:** Perform Smith-Waterman local alignment around each chain to produce the final alignment
- 4. Scoring:** Select the best alignment and assign a MAPQ (mapping quality) score

#### GPU Acceleration Strategy

Parabicks parallelizes the computationally intensive Smith-Waterman extension step across GPU cores. The FM-index lookup (seeding) remains CPU-bound but constitutes a small fraction of total compute. The fq2bam command also integrates coordinate sorting and duplicate marking, eliminating separate samtools sort and picard MarkDuplicates steps.

#### Performance on DGX Spark (GB10)

Metric	Value
Wall time	20-45 minutes
GPU utilization	70-90%

<b>Peak memory</b>	~40 GB (of 128 GB unified)
<b>Output</b>	Sorted BAM + BAI index
<b>Mapping rate</b>	>99.5%
<b>Duplicate rate</b>	~8-12%

## 1.3 Deep Learning Variant Calling: DeepVariant

Google DeepVariant reframes variant calling as an image classification problem. For each candidate variant site, it constructs a pileup image — a visual representation of aligned reads at that position — and classifies it using a convolutional neural network (CNN).

### Architecture Details

**Input:** Pileup image (channels: read bases, base qualities, mapping qualities, strand, etc.)

**Network:** Inception-v3 CNN architecture

**Output:** Three-class softmax (homozygous reference, heterozygous variant, homozygous variant)

**Training:** Supervised on GIAB truth sets, with data augmentation and hard example mining

### Why DeepVariant Outperforms GATK HaplotypeCaller

1. The CNN learns complex error patterns that statistical models cannot capture
2. No explicit error model required — the network learns directly from data
3. Better performance on indels and complex variants
4. Transferable across sequencing platforms (Illumina, PacBio, ONT)

### Performance

Metric	Value
<b>Wall time</b>	10-35 minutes (GPU-accelerated via Parabricks)
<b>GPU utilization</b>	80-95%
<b>Peak memory</b>	~60 GB
<b>SNP F1</b>	>99.7% on HG002
<b>Indel F1</b>	>99.4% on HG002
<b>Total variants</b>	~11.7M (unfiltered)
<b>QUAL&gt;30 variants</b>	~3.5M

## 1.4 VCF Quality Metrics

Metric	Expected Range	Interpretation
--------	----------------	----------------

<b>Ti/Tv ratio</b>	2.0-2.1	Transition/transversion ratio; deviation suggests systematic error
<b>Het/Hom ratio</b>	1.5-2.0	Heterozygous/homozygous ratio; population-dependent
<b>SNP count</b>	~4.2M	Consistent with Ashkenazi ancestry
<b>Indel count</b>	~1.0M	Normal range for WGS
<b>Novel variant rate</b>	<5%	Variants not in dbSNP; higher rates suggest error

# Chapter 2: Variant Annotation — Multi-Database Integration

## 2.1 ClinVar: Clinical Variant Classification

ClinVar (NCBI) is a freely accessible archive of relationships between human variants and phenotypes. The HCLS AI Factory integrates the February 2026 release containing 4.1 million variant-condition records.

### Classification System (ACMG/AMP)

**Pathogenic (P)** — Strong evidence of disease causation

**Likely Pathogenic (LP)** — Moderate evidence

**Variant of Uncertain Significance (VUS)** — Insufficient evidence

**Likely Benign (LB)** — Moderate evidence against pathogenicity

**Benign (B)** — Strong evidence against pathogenicity

### Review Status Tiers

ClinVar classifies assertion confidence using star ratings (0-4 stars). The pipeline weights variants with  $\geq 2$  stars (multiple submitters with concordant interpretations) more heavily.

### Annotation Performance

Of ~3.5M QUAL>30 variants, approximately 35,616 (1.0%) match ClinVar entries. The low match rate reflects that most variants in a healthy individual are common polymorphisms not represented in a clinical database focused on rare disease.

## 2.2 AlphaMissense: AI Pathogenicity Prediction

AlphaMissense (Cheng et al., Science 2023) predicts the pathogenicity of all possible human missense variants using features derived from AlphaFold protein structure predictions and evolutionary conservation.

### Model Architecture

**Input features:** amino acid sequence context, evolutionary conservation (from MSA), and structural features from AlphaFold

**Output:** pathogenicity score (0-1, continuous)

**Total predictions:** 71,697,560 unique missense variants

### Calibrated Thresholds

**Pathogenic:** >0.564 (90% precision on ClinVar pathogenic set)

**Ambiguous:** 0.34-0.564

**Benign:** <0.34 (90% precision on ClinVar benign set)

## Critical Limitation

AlphaMissense only predicts missense variant effects. Stop-gain, frameshift, splice site, and non-coding variants require other prediction tools. The pipeline uses VEP for functional consequence annotation to complement AlphaMissense.

## 2.3 Ensembl VEP: Functional Consequence Prediction

The Variant Effect Predictor maps variants to genes, transcripts, and regulatory regions, annotating each with standardized Sequence Ontology (SO) terms.

### Impact Classification

Impact Level	Example Consequences	Typical Action
HIGH	stop_gained, frameshift_variant, splice_donor_variant	Likely loss of function
MODERATE	missense_variant, inframe_deletion	Protein function may change
LOW	synonymous_variant, splice_region_variant	Unlikely to affect protein
MODIFIER	intron_variant, upstream_gene_variant	Non-coding effects

## 2.4 The Annotation Pipeline Architecture

The three annotation databases are applied sequentially in annotator.py (23 KB):

### Annotation Pipeline Flow

```
VCF (11.7M variants)
→ parse_vcf(min_qual=30)
→ annotate_clinvar()
→ annotate_alphamissense()
→ annotate_vep()
→ generate_text_summary()
→ embed_variants()
→ index_in_milvus()

→ 3.5M variants
→ Clinical significance
→ AI pathogenicity scores
→ Functional consequences
→ Natural language descriptions
→ 384-dim BGE embeddings
→ Searchable vector database
```

# Chapter 3: Vector Database Architecture – Milvus and RAG

## 3.1 Milvus Schema Design

The genomic\_evidence collection in Milvus 2.4 uses a 17-field schema designed to support both vector similarity search and scalar filtering:

Field	Type	Rationale
<code>id</code>	INT64 (PK, auto)	Milvus-managed primary key
<code>embedding</code>	FLOAT_VECTOR(384)	Semantic search vector
<code>chrom</code>	VARCHAR(10)	Genomic coordinate filtering
<code>pos</code>	INT64	Positional queries
<code>ref/alt</code>	VARCHAR(1000)	Allele matching
<code>qual</code>	FLOAT	Quality score filtering
<code>gene</code>	VARCHAR(100)	Gene-level queries
<code>consequence</code>	VARCHAR(200)	Functional filtering (e.g., missense only)
<code>impact</code>	VARCHAR(20)	Impact level filtering
<code>genotype</code>	VARCHAR(10)	Zygosity queries
<code>text_summary</code>	VARCHAR(2000)	Human-readable context for RAG
<code>clinical_significance</code>	VARCHAR(200)	ClinVar classification
<code>rsid</code>	VARCHAR(20)	dbSNP lookup
<code>disease_associations</code>	VARCHAR(2000)	Disease context for RAG
<code>am_pathogenicity</code>	FLOAT	AlphaMissense score filtering
<code>am_class</code>	VARCHAR(20)	Pathogenicity class filtering

## 3.2 Index Configuration and Performance

### Index Type: IVF\_FLAT (Inverted File with Flat Vectors)

**Why IVF\_FLAT?** At 3.5M vectors with 384 dimensions, IVF\_FLAT provides the best recall-latency tradeoff. HNSW would use more memory; IVF\_PQ would sacrifice recall.

**nlist=1024:** Partitions vectors into 1024 clusters. Query searches ~16 clusters (nprobe=16), examining ~55K vectors per query.

**Metric:** COSINE similarity (normalized dot product)

### Search Performance

Metric	Value
<code>Index build time</code>	~8 minutes (3.5M × 384-dim)

<b>Index memory</b>	~2 GB
<b>Search latency (nprobe=16)</b>	8-15 ms
<b>Recall@20</b>	>95%

## 3.3 RAG Architecture with Claude

The RAG pipeline in `rag_engine.py` (23 KB) implements a multi-stage retrieval strategy:

### 1. Query Expansion

User queries are enriched using 10 therapeutic area keyword maps. For example, a query about "neurodegeneration" is expanded with terms like "frontotemporal dementia," "ALS," "motor neuron," "tau protein."

### 2. Hybrid Retrieval

The expanded query is embedded and used for vector search (`top_k=20`). Results are optionally filtered by scalar fields (e.g., `impact=HIGH`, `am_class=pathogenic`).

### 3. Context Assembly

Retrieved variants are formatted into structured context:

#### Context Template

```
## Variant Evidence
- chr9:35065263 G>A | Gene: VCP | Consequence: missense_variant
  ClinVar: Pathogenic | AlphaMissense: 0.87 (pathogenic)
  Disease: Frontotemporal Dementia, ALS, IBMPFD
```

### 4. Claude Inference

The assembled context + knowledge base + user query are sent to `claude-sonnet-4-20250514` (`temperature=0.3`, `max_tokens=4096`).

### Why temperature=0.3?

Lower temperature produces more deterministic, factual responses. For clinical genomics, hallucination is dangerous — the model should report only what the evidence supports.

# Chapter 4: Drug Discovery Pipeline — Deep Dive

## 4.1 The 10-Stage Architecture

The drug discovery pipeline in pipeline.py (18 KB) implements a sequential 10-stage workflow:

Stage	Module	Key Algorithm
1. Initialize	pipeline.py	Pydantic model validation
2. Normalize Target	pipeline.py	Gene → UniProt → PDB mapping
3. Structure Discovery	cryoem_evidence.py	RCSB PDB REST API query
4. Structure Preparation	cryoem_evidence.py	Multi-factor scoring
5. Molecule Generation	nim_clients.py	MolMIM masked LM inference
6. Chemistry QC	molecule_generator.py	RDKit valence/kekulization
7. Conformer Generation	molecule_generator.py	RDKit ETKDG algorithm
8. Molecular Docking	nim_clients.py	DiffDock diffusion inference
9. Composite Ranking	pipeline.py	Weighted multi-objective
10. Reporting	pipeline.py	ReportLab PDF generation

## 4.2 Cryo-EM Structure Scoring

The cryoem\_evidence.py (6 KB) module implements a multi-factor structure scoring algorithm:

### Python

```
score += max(0, 5.0 - resolution)          # Resolution: 0-5 scale
if has_inhibitor_bound: score += 3.0        # Binding site defined
score += num_druggable_pockets * 0.5         # Pocket count bonus
if 'Cryo-EM' in method: score += 0.5          # Method bonus
```

### Design Rationale

**Resolution:** the primary factor (0-5 scale). The 5 Å cutoff excludes low-resolution structures unsuitable for docking.

**Inhibitor bonus (+3):** Inhibitor-bound structures provide a pre-defined binding site and reference ligand geometry.

**Pocket count (+0.5 each):** More druggable pockets increase therapeutic options.

**Cryo-EM bonus (+0.5):** Reflects the growing prevalence and quality of Cryo-EM structures for drug targets.

## 4.3 MolMIM: Molecular Masked Inverse Modeling

MolMIM applies masked language modeling (the technique behind BERT in NLP) to molecular SMILES strings. Given a seed molecule, it:

1. Tokenizes the SMILES into a vocabulary of molecular substructures
2. Randomly masks 15-30% of tokens
3. Predicts the masked tokens using a transformer architecture
4. The predicted tokens create novel molecular structures

## Critical Considerations

**SMILES output:** MolMIM generates SMILES strings, not 3D structures. Chemical validity must be verified by RDKit.

**Stochastic generation:** Different random seeds produce different molecules.

**Temperature control:** Higher temperature = more diverse but potentially less valid molecules.

## 4.4 DiffDock: Diffusion-Based Molecular Docking

DiffDock (Corso et al., ICLR 2023) models molecular docking as a generative diffusion process over the product space of rotations, translations, and torsion angles.

### Key Innovation

Unlike grid-based docking methods (AutoDock Vina, Glide), DiffDock does not require a pre-defined search box around a binding site. It learns to predict binding poses directly from protein-ligand pairs, making it suitable for blind docking.

### Score Interpretation

**Confidence score (0-1):** indicates the model's certainty about the predicted pose

**Binding affinity (kcal/mol):** estimates the free energy of binding; more negative = stronger binding

### Limitations

**Training bias:** DiffDock was trained primarily on crystal structures; performance may degrade on Cryo-EM structures with lower resolution

**No kinetics:** The model predicts pose and affinity but not binding kinetics (on/off rates)

**Rigid protein:** Protein flexibility is not modeled — the protein is treated as rigid

## 4.5 Composite Scoring and Normalization

The composite scoring formula balances three objectives:

### Python

```
dock_normalized = max(0.0, min(1.0, (10.0 + dock_score) / 20.0))
composite = 0.30 * gen_score + 0.40 * dock_normalized + 0.30 * qed_score
```

### Normalization Rationale

**Docking scores:** range from ~-15 to ~0 kcal/mol. The formula  $(10 + \text{dock}) / 20$  maps this to approximately 0-1, with -10 kcal/mol mapping to 0.0 and +10 mapping to 1.0.

**Generation scores:** already 0-1 (MolMIM confidence).

**QED scores:** inherently 0-1.

## Weight Rationale

**Docking (40%):** receives the highest weight because binding affinity is the most direct predictor of therapeutic activity

**Generation (30%):** balances novelty of the molecular design

**QED (30%):** balances practical drug-likeness

# Chapter 5: Nextflow DSL2 Pipeline Architecture

## 5.1 Module Design

The pipeline uses Nextflow DSL2's module system for composable workflow design:

### Directory Structure

```
hls-orchestrator/
├── main.nf          # Entry point, mode routing
├── nextflow.config   # Profiles, parameters
├── run_pipeline.py   # Python CLI launcher
└── modules/
    ├── genomics.nf    # Stage 1 processes
    ├── rag_chat.nf    # Stage 2 processes
    ├── drug_discovery.nf # Stage 3 processes
    └── reporting.nf   # Report generation
```

## 5.2 Execution Modes and Data Flow

Mode	Data Flow	Use Case
<b>full</b>	FASTQ → VCF → Target → Candidates	Complete pipeline
<b>target</b>	VCF → Target → Candidates	Pre-existing VCF
<b>drug</b>	Target → Candidates	Known gene target
<b>demo</b>	Pre-configured FASTQ → Candidates	VCP/FTD demonstration
<b>genomics_only</b>	FASTQ → VCF	Variant calling only

## 5.3 Profile Configuration

The `nextflow.config` defines six execution profiles optimized for different environments:

**dgx\_spark:** GPU resource requests, memory limits tuned for 128 GB unified memory

**docker:** Docker container execution with GPU passthrough

**singularity:** Singularity containers for HPC environments without Docker

**slurm:** SLURM scheduler integration for cluster execution

# Chapter 6: Clinical Translation and Limitations

## 6.1 From Computational Hits to Drug Leads

The HCLS AI Factory generates computational drug candidates — not approved medications. The path from computational hit to clinical drug requires:

- 1. In vitro validation:** Test top candidates in biochemical assays (e.g., VCP ATPase activity inhibition)
- 2. Cell-based assays:** Confirm activity in relevant cell lines
- 3. ADMET profiling:** Absorption, Distribution, Metabolism, Excretion, and Toxicity studies
- 4. Lead optimization:** Iterative cycles of design, synthesis, and testing
- 5. Preclinical studies:** Animal models for efficacy and safety
- 6. Clinical trials:** Phase I (safety), Phase II (efficacy), Phase III (large-scale)

### Estimated Timeline

10-15 years from computational hit to approved drug. The HCLS AI Factory accelerates the earliest stage — computational lead generation — from months to minutes.

## 6.2 Limitations and Caveats

### Genomics

- DeepVariant accuracy varies by variant type (SNPs > indels > structural variants)
- Short-read WGS has limited sensitivity for structural variants and repeat expansions
- Population-specific biases in GRCh38 may affect variant calling in non-European ancestries

### RAG/Annotation

- ClinVar has known biases toward well-studied genes and European ancestry variants
- AlphaMissense is limited to missense variants; non-coding variants are not scored
- The 201-gene knowledge base covers common drug targets but not the full druggable genome

### Drug Discovery

- MolMIM-generated molecules have not been synthesized or tested
- DiffDock docking scores are predictions, not experimental measurements
- Protein flexibility is not modeled; induced-fit effects are ignored
- The composite scoring weights (30/40/30) are heuristic, not optimized on clinical outcomes

## 6.3 Ethical Considerations

**Informed consent:** Patient genomic data requires explicit consent for research use

**Data sovereignty:** NVIDIA FLARE federated learning keeps data local; essential for HIPAA/GDPR compliance

**Return of results:** Incidental findings (e.g., BRCA1 pathogenic variants) may require clinical reporting

**Equity:** Pipeline performance should be validated across diverse ancestries to avoid exacerbating health disparities

# Chapter 7: Scaling Analysis

## 7.1 DGX Spark Bottleneck Analysis

Component	Bottleneck	Phase 1 Impact
Parabricks (fq2bam)	GPU compute	20-45 min, acceptable
DeepVariant	GPU memory (60 GB peak)	Leaves 68 GB for other tasks
Milvus indexing	CPU + I/O	24 min for 3.5M vectors
MolMIM inference	GPU compute	2 min for 100 molecules
DiffDock inference	GPU compute + memory	8 min for 98 candidates
Sequential total	GPU time-sharing	~4 hours end-to-end

## 7.2 Phase 2: DGX B200 Scaling

With 8x B200 GPUs and 1-2 TB HBM3e:

**Parallel Parabricks:** 4-8 simultaneous samples

**Dedicated Milvus GPU:** GPU-accelerated vector search (sub-millisecond)

**NIM replicas:** 2-4 MolMIM + 2-4 DiffDock instances

**Estimated throughput:** 10-20 patients per day

## 7.3 Phase 3: DGX SuperPOD

**Hundreds of B200 GPUs** with NVLink and InfiniBand

**Distributed Milvus cluster:** Billions of variants across institutions

**NVIDIA FLARE:** Federated model training without data sharing

**Estimated throughput:** Thousands of patients per day

# Chapter 8: Advanced Topics and Extensions

## 8.1 Alternative Embedding Strategies

BGE-small-en-v1.5 (384-dim) was chosen for its balance of quality and efficiency.  
Alternatives:

Model	Dimensions	Size	Trade-off
<b>BGE-small-en-v1.5</b>	384	33M params	Current choice: fast, efficient
<b>BGE-base-en-v1.5</b>	768	109M params	Better recall, 2× memory
<b>BGE-large-en-v1.5</b>	1024	335M params	Best recall, 3× memory
<b>BiomedBERT</b>	768	109M params	Domain-specific, biomedical text
<b>PubMedBERT</b>	768	109M params	PubMed-trained, clinical text

## 8.2 Multi-Objective Optimization

The current composite scoring uses fixed weights (30/40/30). Advanced approaches:

**Pareto optimization:** Identify the Pareto frontier of generation, docking, and QED

**Bayesian optimization:** Learn optimal weights from experimental feedback

**Active learning:** Prioritize candidates that reduce uncertainty in the scoring model

## 8.3 Long-Read Sequencing Integration

Oxford Nanopore and PacBio long-read technologies can detect structural variants (SVs) and repeat expansions that short-read WGS misses. Future extensions could:

- Add ONT/PacBio alignment with minimap2
- Detect SVs with Sniffles2 or PEPPER-Margin-DeepVariant
- Phase haplotypes for compound heterozygosity detection

## 8.4 Pharmacogenomics Integration

The knowledge base includes 11 pharmacogenomics genes (CYP2D6, CYP2C19, CYP3A4, DYPD, TPMT, etc.). Future extensions could:

- Star allele calling with PharmCAT
- Drug-drug interaction prediction
- Dosing recommendations based on metabolizer status



# Review Questions

Graduate-level questions for self-assessment and discussion. These require synthesis across multiple chapters.

- 1.** Explain why DeepVariant reframes variant calling as an image classification problem. What advantages does this provide over statistical models like GATK HaplotypeCaller?
- 2.** Describe the IVF\_FLAT index configuration (`nlist=1024, nprobe=16`) and calculate the approximate number of vectors examined per query from a collection of 3.5M vectors.
- 3.** Why does the RAG pipeline use `temperature=0.3` for Claude? What are the trade-offs of lower vs. higher temperature in clinical genomics applications?
- 4.** Explain the docking score normalization formula  $\max(0, \min(1, (10 + \text{dock\_score}) / 20))$ . What docking score maps to 0.5? Why is this mapping appropriate?
- 5.** Compare the AlphaMissense thresholds (`pathogenic >0.564, benign <0.34`) with ClinVar classifications. What does the "ambiguous" zone represent, and why is it clinically significant?
- 6.** Describe three limitations of DiffDock that could affect the reliability of the drug candidate rankings.
- 7.** Explain why inhibitor-bound PDB structures (like 5FTK) receive a +3 bonus in the structure scoring algorithm. What information does an inhibitor-bound structure provide that an apo structure does not?
- 8.** Design a validation experiment to test the top 5 drug candidates from the VCP/FTD demo pipeline. What assays would you use, and what would constitute a positive result?
- 9.** Calculate the approximate memory budget for the Milvus vector index: 3.5M vectors  $\times$  384 dimensions  $\times$  4 bytes per float. How does this compare to the available memory on DGX Spark?
- 10.** The composite scoring weights (30% generation, 40% docking, 30% QED) are heuristic. Propose an approach to optimize these weights using experimental feedback from in vitro screening.